

Module 2 – Part I

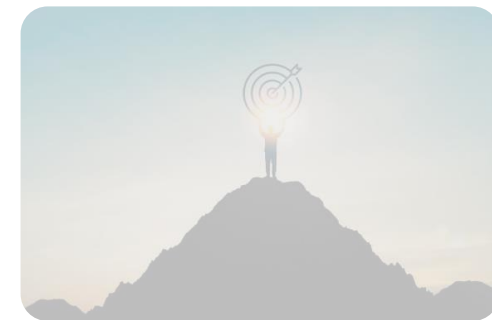
Setting up Deep Forecasting Environment (Python)





Road map!

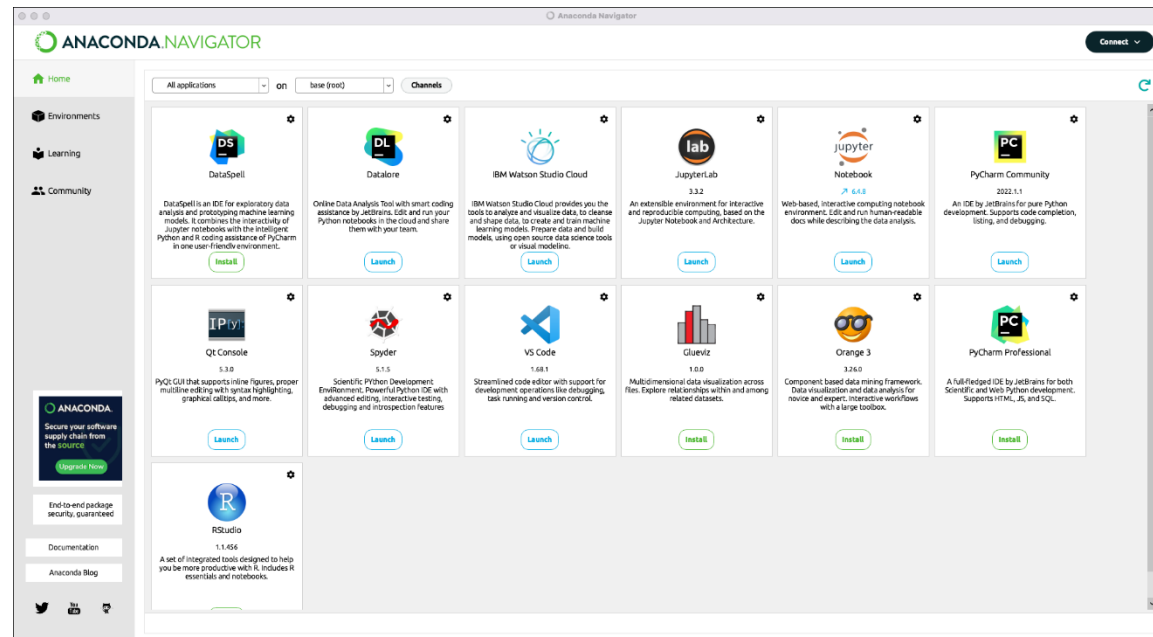
- Module 1- Introduction to Deep Forecasting
- **Module 2- Setting up Deep Forecasting Environment (Python)**
- Module 3- Exponential Smoothing
- Module 4- ARIMA models
- Module 5- Machine Learning for Time series Forecasting
- Module 6- Deep Neural Networks
- Module 7- Deep Sequence Modeling (RNN, LSTM)
- Module 8- Transformers (Attention is all you need!)
- Module 9- Prophet and Neural Prophet

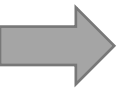


Install through



- Anaconda is a **distribution** of the Python and R programming languages for scientific computing, that aims to simplify package management with conda **environments**.
- Anaconda offers the easiest way to perform data science and machine learning on a single machine.
- Install Anaconda @ <https://www.anaconda.com/>

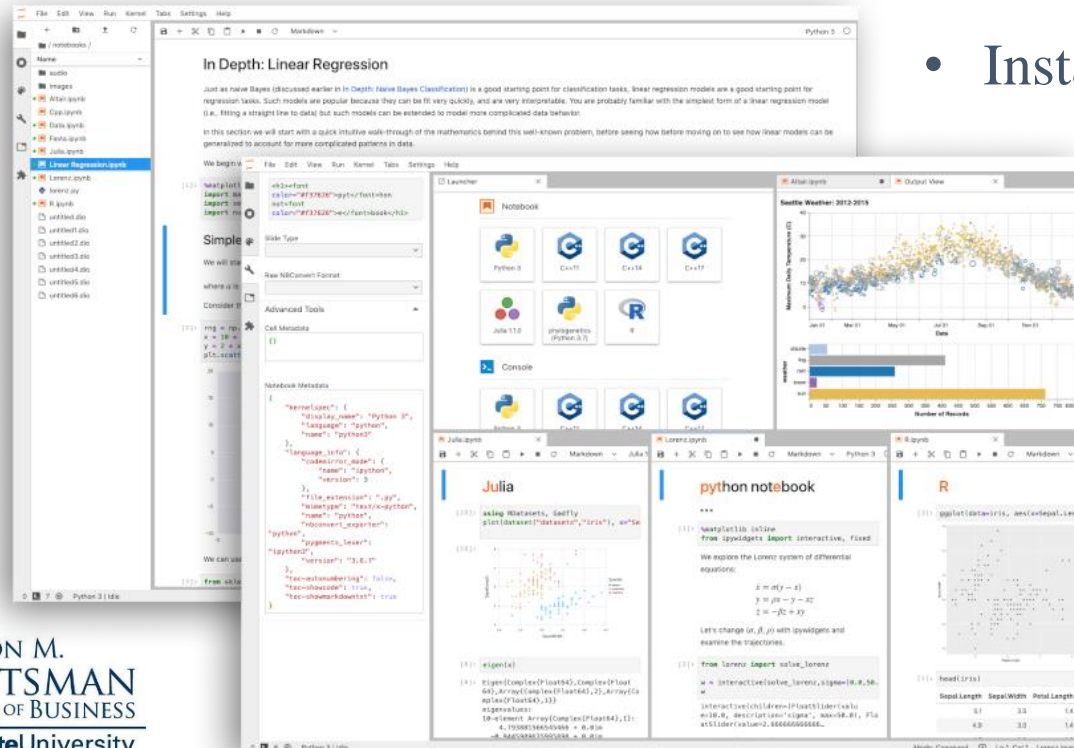




JupyterLab



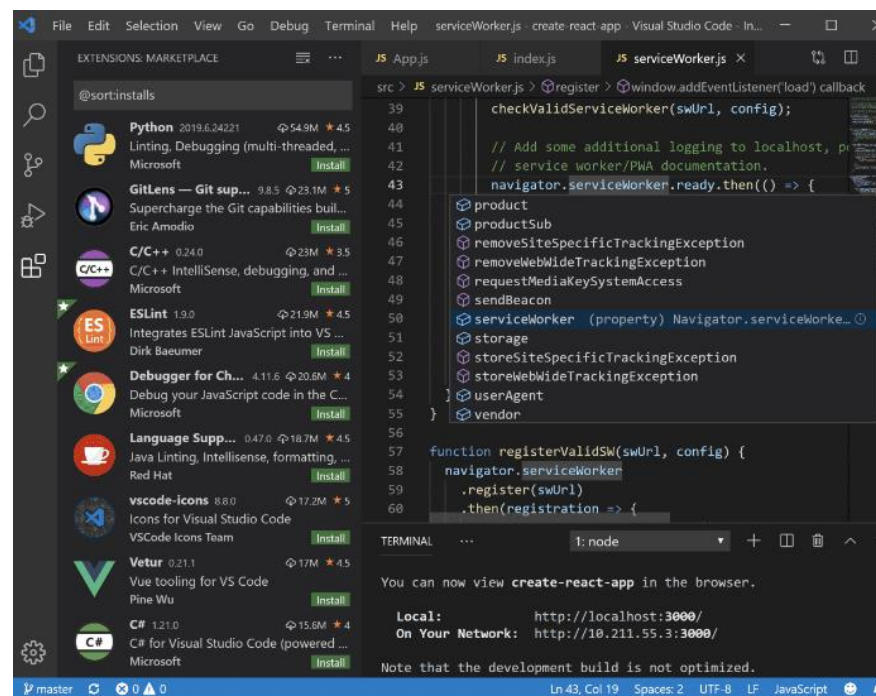
- JupyterLab is the latest **web-based interactive development environment** for notebooks, code, and data
- Jupyter's name is a reference to the three core programming languages supported by Jupyter, which are **Julia**, **Python** and **R**



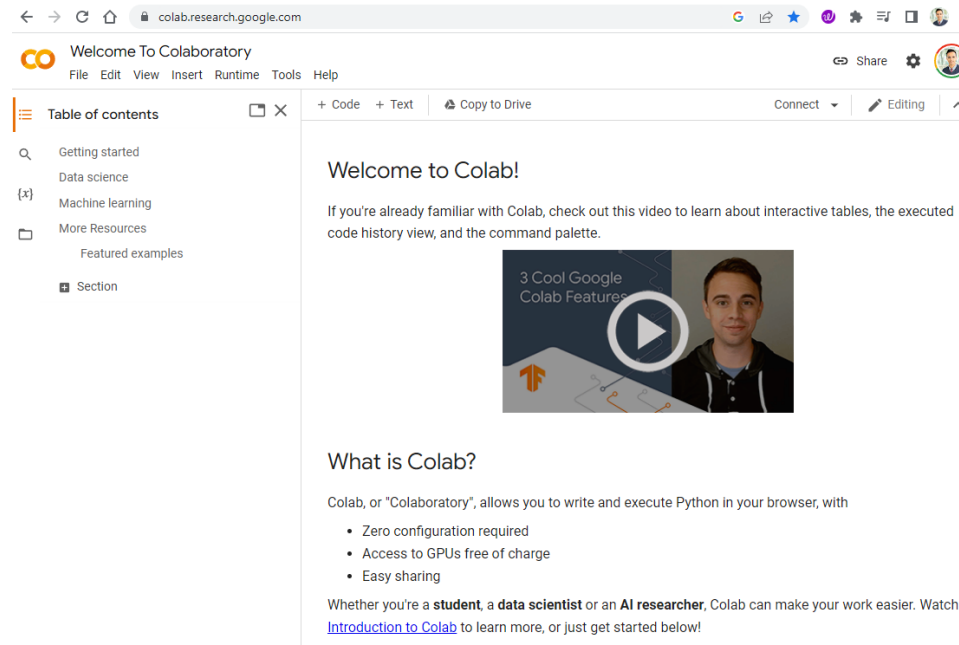
- Install JupyterLab @ <http://jupyter.org/install>

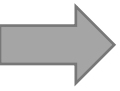
VS Code

- VS Code is one of the most popular source code editors
- Features include support for **debugging**, syntax highlighting, intelligent **code completion**, code refactoring, and **embedded Git**.
- Install VS code @ <https://code.visualstudio.com/>



- Colab is a free hosted **Jupyter notebook-style environment** that runs entirely in the **cloud** and requires no setup to use. It also provides access to **machine learning libraries** and computing resources including **GPU**.
- Colab allows anybody to write and execute arbitrary **python code** through the **browser**, and is especially well suited to machine learning, data analysis and education. <https://colab.research.google.com/>



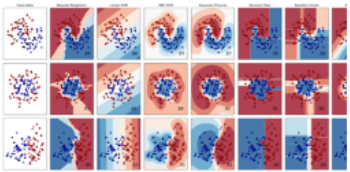


Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.

Algorithms: SVM, nearest neighbors, random forest, and more...



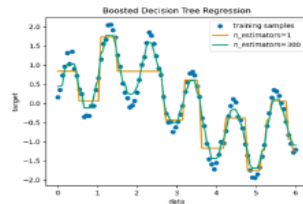
Examples

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: SVR, nearest neighbors, random forest, and more...



Examples

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, and more...



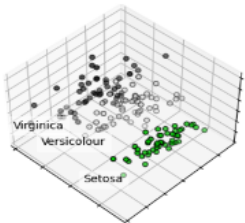
Examples

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, feature selection, non-negative matrix factorization, and more...



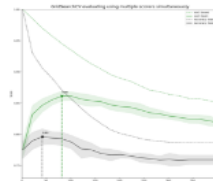
Examples

Model selection

Comparing, validating and choosing parameters and models.

Applications: Improved accuracy via parameter tuning

Algorithms: grid search, cross validation, metrics, and more...



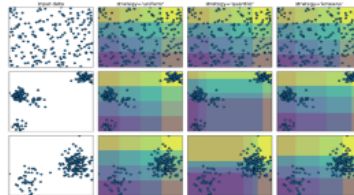
Examples

Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text for use with machine learning algorithms.

Algorithms: preprocessing, feature extraction, and more...



Examples

- Scikit-learn is an **open-sourced Python** library and includes a variety of unsupervised and supervised learning techniques.
- It is based on technologies and libraries like Matplotlib, Pandas and NumPy and helps simplify the coding task.
- Install Scikit-learn @ <https://scikit-learn.org/stable/install.html>



- PyCaret is an **open-source**, **low-code** machine learning library in Python that automates machine learning workflows.
- PyCaret is essentially a **Python wrapper** around several machine learning libraries and frameworks
- Install PyCaret @ <https://pycaret.gitbook.io/docs/get-started/installation>

```
# load dataset
import pandas as pd
train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')

# init setup
from pycaret.classification import *
s = setup(train, target= 'target')

# model training and selection
best = compare_models()

# analyze best model
evaluate_model(best)

# predict on new data
predictions = predict_model(best, data =test )

# save best pipeline
save_model(best, 'my_best_pipeline')
```

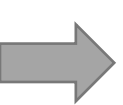
ID	Name
'lr'	Logistic Regression
'knn'	K Nearest Neighbour
'nb'	Naives Bayes
'dt'	Decision Tree Classifier
'svm'	SVM - Linear Kernel
'rbfsvm'	SVM - Radial Kernel
'gpc'	Gaussian Process Classifier
'mlp'	Multi Level Perceptron
'ridge'	Ridge Classifier
'rf'	Random Forest Classifier
'qda'	Quadratic Discriminant Analysis
'ada'	Ada Boost Classifier
'gbc'	Gradient Boosting Classifier
'lda'	Linear Discriminant Analysis
'et'	Extra Trees Classifier
'xgboost'	Extreme Gradient Boosting
'lightgbm'	Light Gradient Boosting
'catboost'	CatBoost Classifier



K Keras

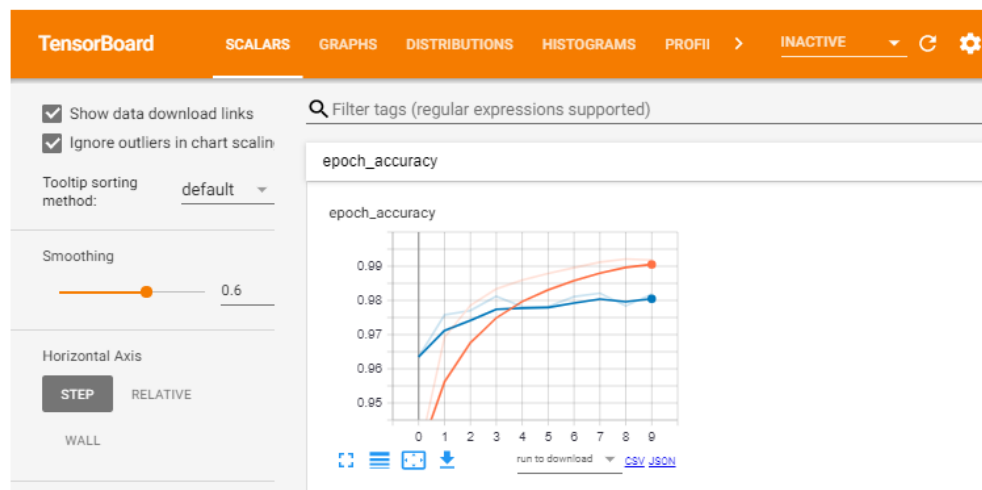
- Keras is a **high-level**, open-source **neural network** library written in Python. It was developed to make it easier for researchers and developers to build and experiment with deep learning models.
- The Keras API became the official high-level API for TensorFlow 2.0 in **2019**.
<https://keras.io/>

```
model = keras.Sequential(  
    [  
        keras.Input(shape=input_shape),  
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Flatten(),  
        layers.Dropout(0.5),  
        layers.Dense(num_classes, activation="softmax"),  
    ]  
)  
  
model.summary()
```



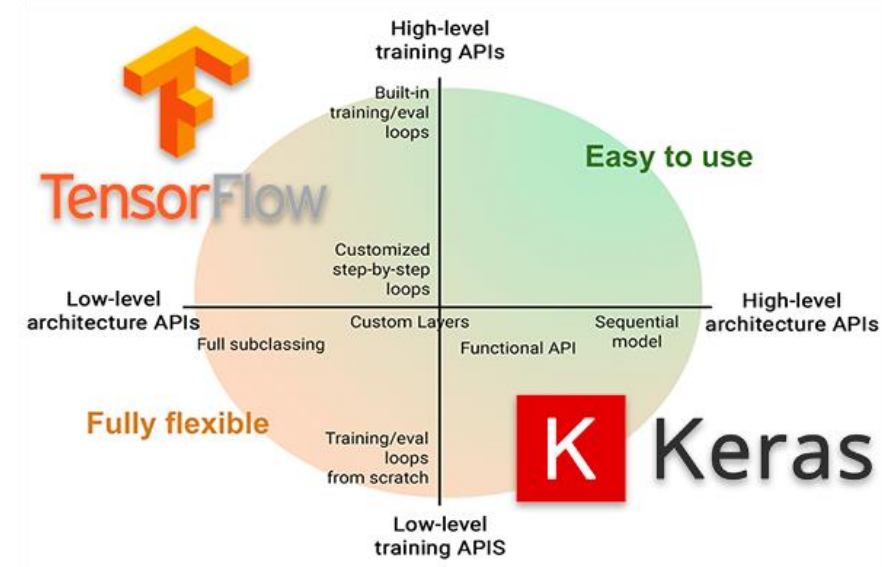
TensorFlow

- TensorFlow is a Google-maintained open-source end-to-end platform for prototyping and assessing machine learning models, **primarily neural networks**.
- TensorFlow also offers **TensorBoard**, a visualization tool for comparing and tracking our learned models.
- It can scale from a single CPU, to a GPU or cluster of GPUs all the way up to a multi-node TPU infrastructure.
- Build-in Google Colab. For local installation visit <https://www.tensorflow.org/install>

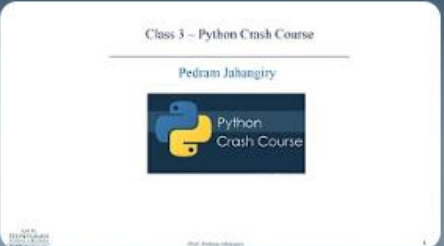


➔ K Keras Vs TF TensorFlow

- **Level of abstraction:** Keras is a **higher-level** library that provides a more intuitive interface for building and training models, while TensorFlow is a **lower-level** library that provides more flexibility but requires the user to specify more details of the model.
- Keras is a **standalone** library, while TensorFlow includes both a low-level library for numerical computations and a high-level library for building and training machine learning models.
- Keras is a user-friendly interface to TensorFlow



Available YouTube playlists



Python Crash Course


Pedram Jahangiry

Public

9 videos 1,612 views Updated today

[Play all](#) [Shuffle](#)

Codes are available on my GitHub account:
<https://github.com/PJalgotrader/platforms-and-tools>



Programming tips


Pedram Jahangiry

Public

13 videos 194 views Updated 3 days ago

[Play all](#) [Shuffle](#)

Codes are available on my GitHub account:
<https://github.com/PJalgotrader/platforms-and-tools>



Google Colab


Pedram Jahangiry

Public

3 videos 73 views Updated 6 days ago

[Play all](#) [Shuffle](#)

Codes are available on my GitHub account:
<https://github.com/PJalgotrader/platforms-and-tools>



PyCaret (Automated machine learning Python package)

Pedram Jahangiry

Public

3 videos 26 views Updated 3 days ago

[Play all](#) [Shuffle](#)

All you need from PyCaret Python library to automate your ML workflow.
Codes are available on my GitHub account:
<https://github.com/PJalgotrader/platforms-and-tools>

→ Platforms and Packages

Listed below are some Python packages and platforms that will be used in the deep learning and deep forecasting courses.




General Python libraries



Machine Learning libraries



Deep Learning libraries





Setting up Deep Learning Environment



Personal Workstation



Cloud Platforms



Google Colaboratory

Pros

- Full control over hardware and software
- Work offline
- Fixed cost

Cons

- Scalability
- Maintenance (both hardware and software)

- Powerful computing resources
- Scalability
- Ease of use
- Cost-effective: Pay-as-you-go
- Collaboration

- Expensive for large-scale experiments
- Dependency on the provider
- Limited control
- Internet connection
- Security

- Powerful computing resources (GPU, TPU)
- Ease of use
- Collaboration
- No need to set up a local environment

- Time limit
- Hardware limitation
- Data storage
- Limited control
- Internet connection
- Security



The modern machine learning landscape

- From 2016 to 2020, the entire machine learning and data science industry has been dominated by these **two approaches**:
 1. Deep learning
 2. Gradient boosted trees
- Most practitioners of deep learning use **Keras**, often in combination with its parent framework **TensorFlow**.
- This means you'll need to be familiar with **Scikit-learn**, **XGBoost**, and **Keras**

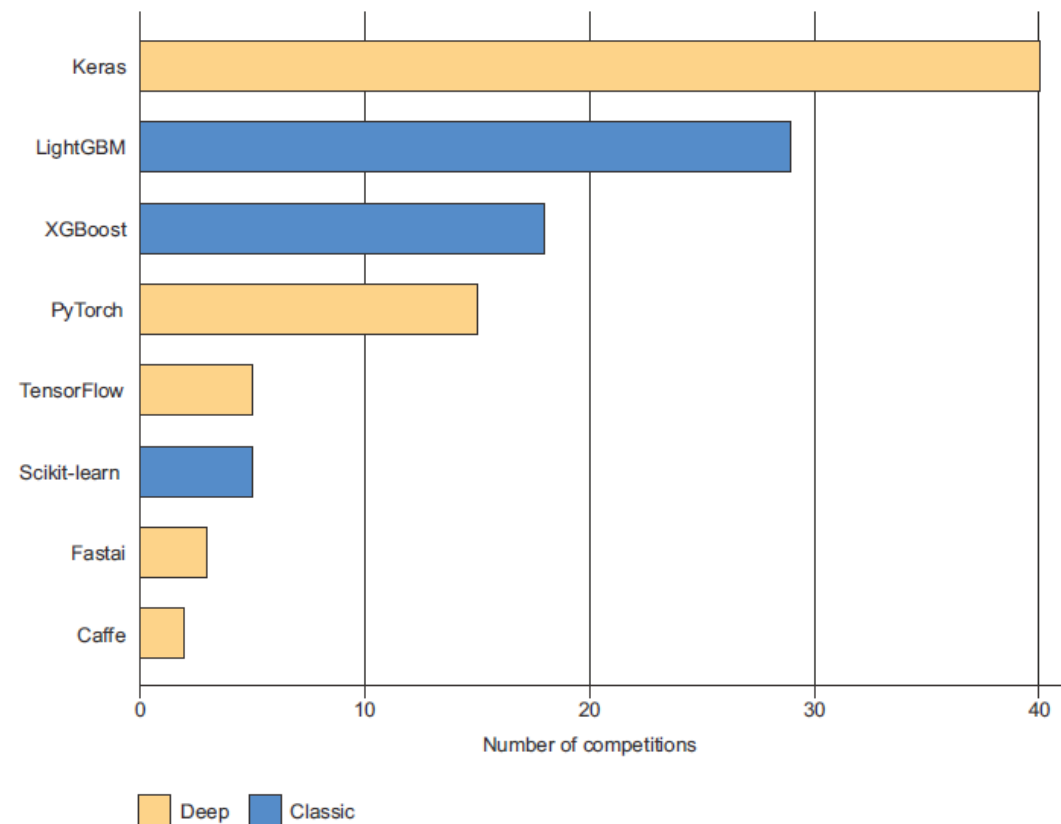


Figure 1.12 Machine learning tools used by top teams on Kaggle



The modern machine learning landscape




Figure 1.13 Tool usage across the machine learning and data science industry (Source: www.kaggle.com/kaggle-survey-2020)

- Kaggle also runs a yearly survey among machine learning and data science professionals worldwide.
- This survey is one of our **most reliable** sources about the **state of the industry!!!**
- This figure shows the percentage of **usage of** different machine learning software frameworks.


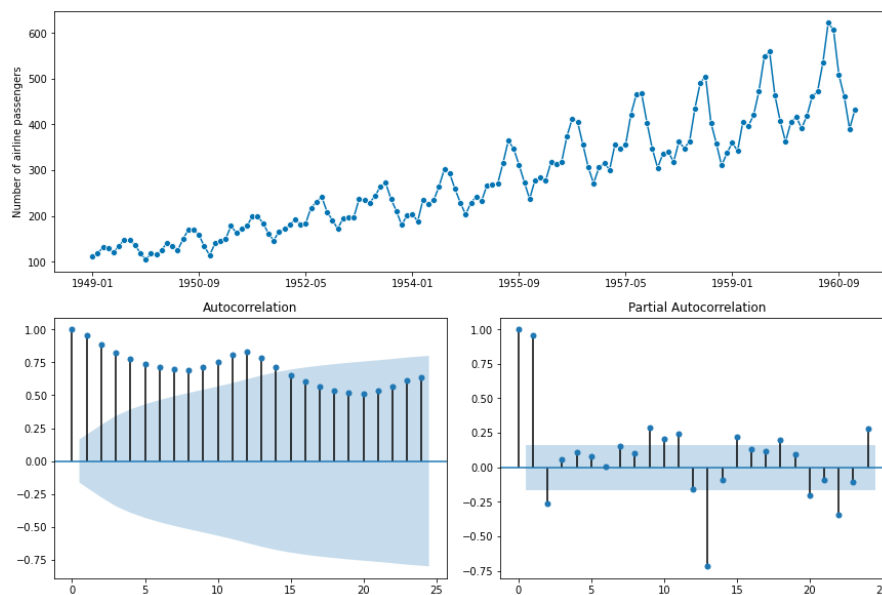


Pedram Jahangiry



Module 2 – Part II


Setting up Deep Forecasting Environment (Time Series Basics)





Road map!

- Module 1- Introduction to Deep Forecasting
- **Module 2- Setting up Deep Forecasting Environment (Time Series Basics)**
- Module 3- Exponential Smoothing
- Module 4- ARIMA models
- Module 5- Machine Learning for Time series Forecasting
- Module 6- Deep Neural Networks
- Module 7- Deep Sequence Modeling (RNN, LSTM)
- Module 8- Transformers (Attention is all you need!)
- Module 9- Prophet and Neural Prophet




→ Time Series Patterns

- **Trend:** A trend is a **long-term** movement in the data, either **upwards** or **downwards**. A trend can be **linear**, meaning that it follows a straight line, or it can be **nonlinear**, meaning that it follows a more complex curve.



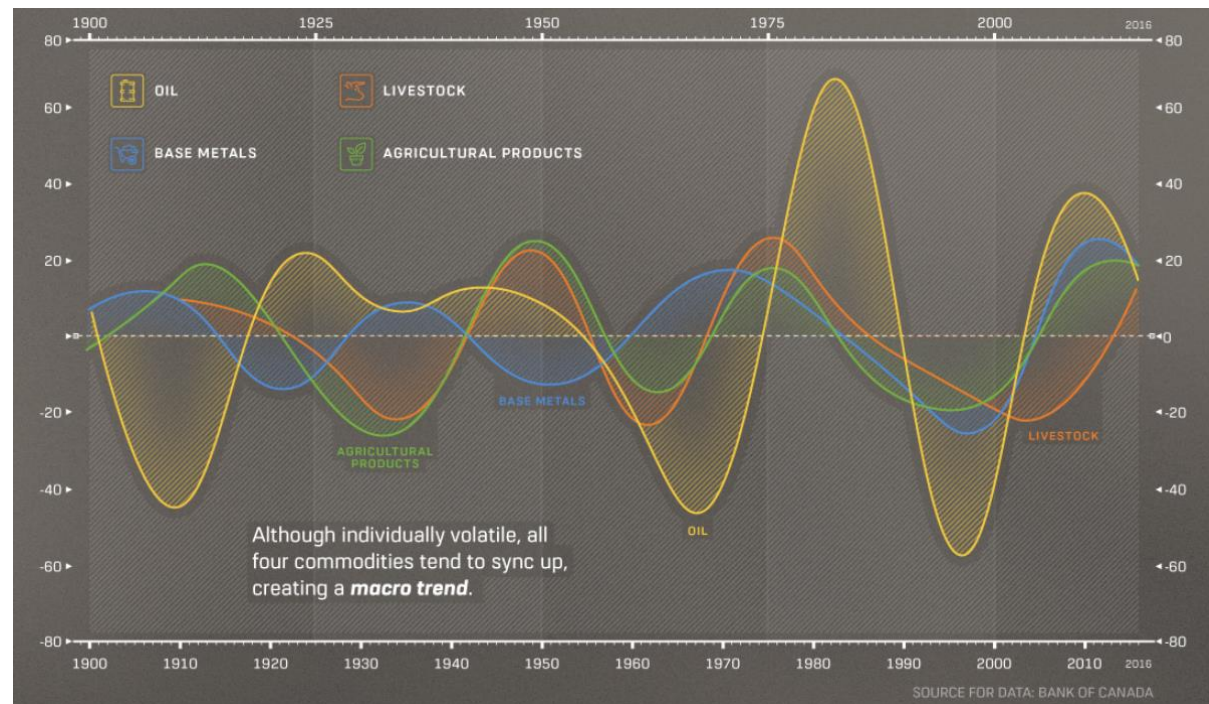
→ Time Series Patterns

- **Seasonal:** A seasonal pattern is a **regular fluctuation** in the data that occurs at a **specific time** of the year, month or week. Seasonality is always of a fixed and known period.



Time Series Patterns

- **Cyclic:** There is a cycle when the data exhibits **peaks** and **valleys** that **do not occur at a fixed frequency**. Economic conditions and the "business cycle" are typically responsible for these fluctuations.
- In general, the **average length** of cycles is **longer** than the length of a seasonal pattern, and the **magnitudes** of cycles tend to be **more variable** than the magnitudes of seasonal patterns.





Lag variables

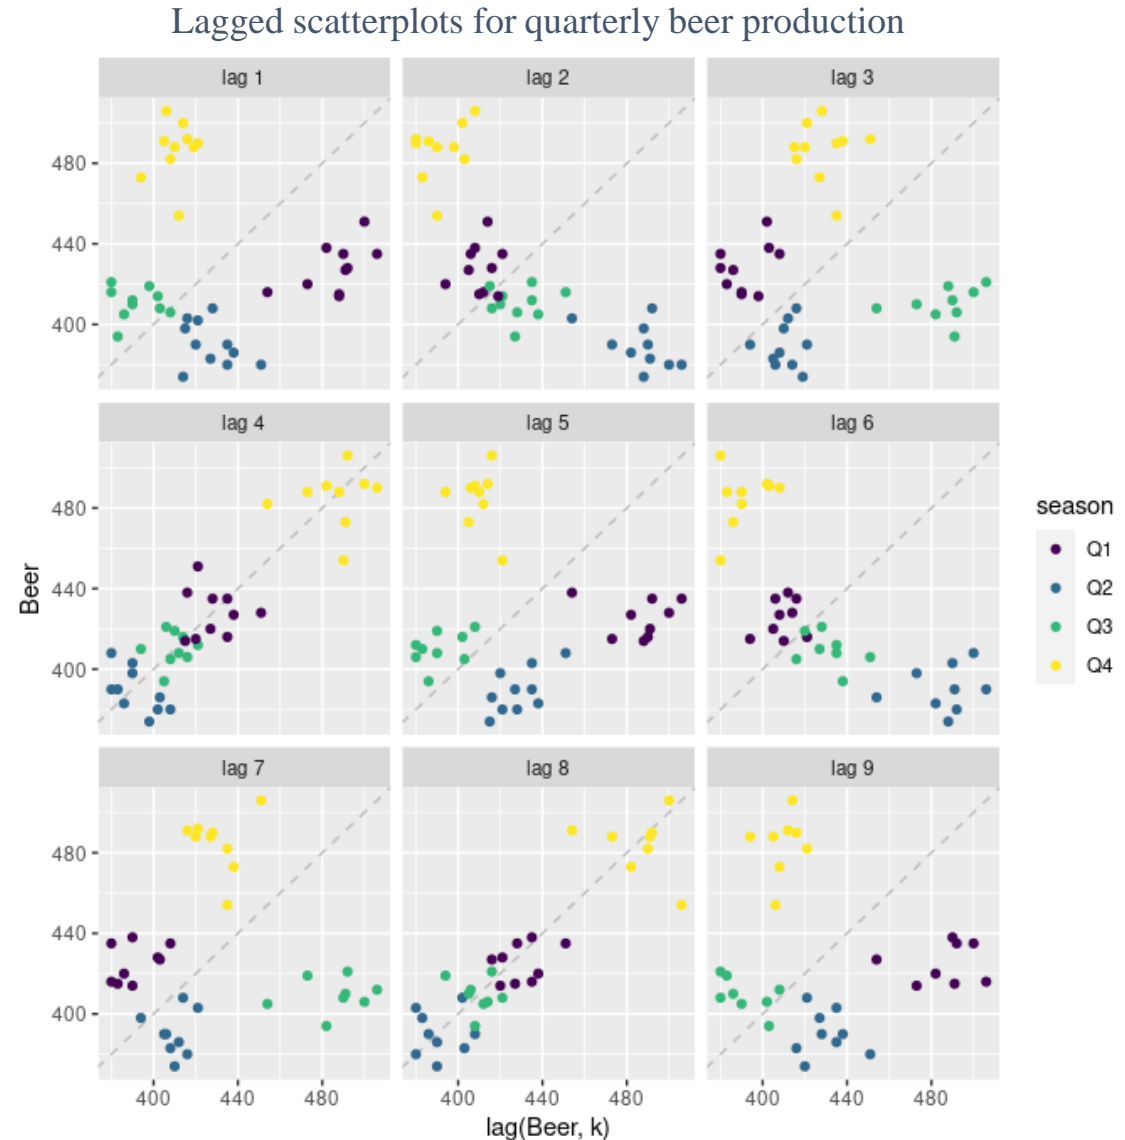
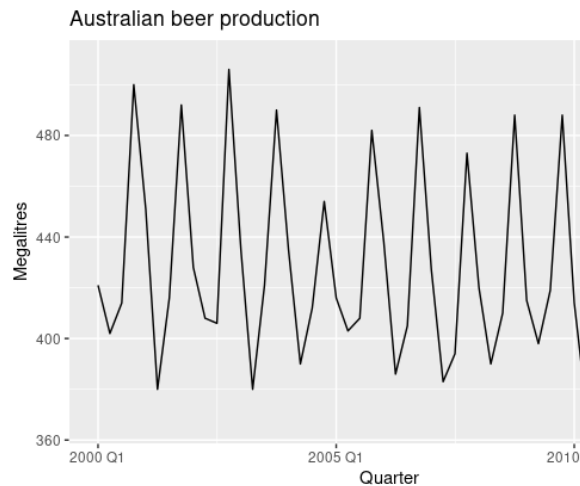
- Lag variables are variables that **are measured at a previous point in time relative to the current time** period being studied.
- Lag variables can be useful for **identifying trends, patterns, and relationships** in time series data and **how they change over time**.

Date	Value	Value _{t-1}	Value _{t-2}
1/1/2017	200	NA	NA
1/2/2017	220	200	NA
1/3/2017	215	220	200
1/4/2017	230	215	220
1/5/2017	235	230	215
1/6/2017	225	235	230
1/7/2017	220	225	235
1/8/2017	225	220	225
1/9/2017	240	225	220
1/10/2017	245	240	225



Lag plots

- y_t plotted against $y_{(t-k)}$
- The colors indicate the quarter of the beer production
- The relationship is **strongly positive** at lags 4 and 8, reflecting the strong **seasonality** in the data.





Autocorrelation

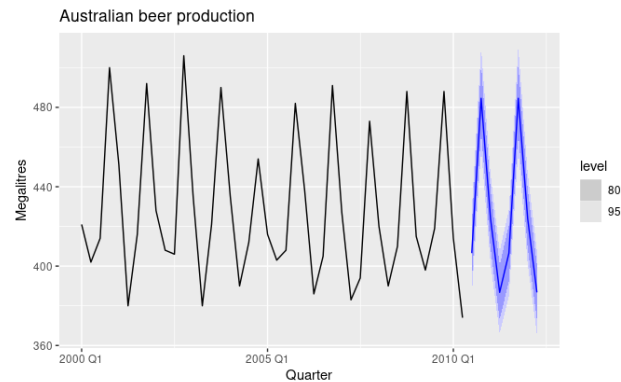
- **Autocorrelation**, also known as **serial correlation**, is a measure of the correlation between a time series and a lagged version of itself.
- It is used to assess the **degree to which** the past values of a time series **are predictive** of its future values.

$$r_k = \frac{\sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})}{\sum_{t=1}^T (y_t - \bar{y})^2}$$




ACF: Autocorrelation Function

- The autocorrelation function (**ACF**) is a statistical tool that can be used to measure the autocorrelation of a time series.
- It calculates the correlation between the time series and lagged versions of itself at different lag periods.





r_1	r_2	r_3	r_4	r_5	r_6	r_7	r_8	r_9
-0.102	-0.657	-0.060	0.869	-0.089	-0.635	-0.054	0.832	-0.108





Trend and seasonality in ACF plots


- Autocorrelation can be useful for **identifying patterns and trends** in time series data.
- The ACF of **trended** time series tend to have **positive values that slowly decrease** as the lags increase.
- For seasonal data, the autocorrelations **are larger** for the **seasonal lags** than for other lags.

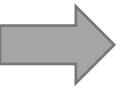


➔ Partial Autocorrelation

- **Partial autocorrelation**, also known as **partial serial correlation**, is a measure of the correlation between a time series and a lagged version of itself, **controlling for the effects of intermediate lag periods**.
- y_t and y_{t-2} might be correlated, simply because they are both connected to y_{t-1} , rather than because of any new information contained in y_{t-2} . **Partial autocorrelation** overcomes this problem.


US Annual GDP per capita (1960-2017)





PACF: Partial Autocorrelation Function



- PACF is a statistical tool that can be used to measure the partial autocorrelation of a time series.

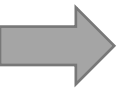




White Noise

- White noise can be thought of as a random **sequence of iid values** (independent and identically distributed) characterized by a distribution.
- By comparing the characteristics of the time series data to the characteristics of white noise, we can determine **whether there are any significant patterns or trends** present in the data.
- White noise data show **no autocorrelation**.





Transformations

- The purpose of **transformations** is to **simplify** the patterns in the historical data by **making the pattern more consistent** across the whole data set.
- Mathematical Transformations:
 - Logarithmic
 - Power
 - Box-Cox
- Note that these transformations should be used with caution, as they can **significantly alter the interpretation** of the data.
- It's generally a good idea to try multiple transformations and compare the results to determine which is most appropriate for your data.





Transformations: Log vs Power vs Box-Cox

Method	Description	Formula
Logarithmic	<ul style="list-style-type: none">Transforms data by taking the logarithm of the values.Log transformation is very interpretable!	$w_t = \log(y_t)$
Power	<ul style="list-style-type: none">Transforms data by raising it to a power. (square, cube, ...)Not so interpretable.	$w_t = y_t^p$
Box-Cox	<ul style="list-style-type: none">Includes both log and power transformations.Depends on parameter lambda.The log is always natural log.	$w_t = \begin{cases} \log(y_t) & \text{if } \lambda = 0; \\ (\text{sign}(y_t) y_t ^\lambda - 1)/\lambda & \text{otherwise.} \end{cases}$


- Common use case: Reducing the impact of **outliers** and making patterns in the data **more visible**.




Box-Cox transformation and lambda

- A good value of λ is one which makes the size of the seasonal variation about the same across the whole series

Australian Quarterly Gas Production



Australian Quarterly Gas Production




→ Time Series Components

- Time series patterns: Trend, Seasonality, Cycle

- **Decomposing** Time Series:

- **Trend-Cycle** component
- **Seasonal** component
- **Remainder** (error) component



- The decomposition of a time series is done to **improve understanding**, but it can also be used to **improve forecasting accuracy**.
- For simplified decomposition (and later analysis), it is sometimes **helpful to transform** or adjust the series first.



Time Series Components Decomposition

TS Decomposition

Additive


$$y_t = S_t + T_t + R_t$$

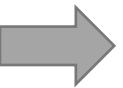
Multiplicative

$$y_t = S_t * T_t * R_t$$


$$\log y_t = \log S_t + \log T_t + \log R_t$$

- The additive decomposition is appropriate if the **magnitude of seasonal fluctuations or variation around trend-cycles do not change with time** series level.






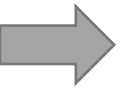
Time Series Components Decomposition Example



➔ Seasonally Adjusted Data

- If the seasonal component is removed from the original data, the resulting values are the “seasonally adjusted” data.



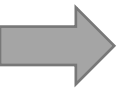


Some Forecasting Benchmarks


- Some forecasting methods are **extremely simple** and **surprisingly effective**.

1. **Mean** method: the forecasts of all future values are equal to **the average of the historical data**.
2. **Naïve** method: all forecasts are set to be the **value of the last observation**
3. **Seasonal Naïve**: each forecast set to be equal to the last observed value **from the same season**
4. **Drift** method: allows forecasts to increase or decrease over time, where the **drift** is set to be the average change in historical data.

- These methods will serve as benchmarks rather than the method of choice.
- If your model cannot beat the benchmark, it is not worth considering!!!



Some Forecasting Benchmarks (Google Stock price forecasting)



➔ Road map!

- ✓ Module 1- Introduction to Deep Forecasting
- ✓ Module 2- Setting up Deep Forecasting Environment
- Module 3- ETS and Exponential Smoothing
- Module 4- ARIMA models
- Module 5- Machine Learning for Time series Forecasting
- Module 6- Deep Neural Networks
- Module 7- Deep Sequence Modeling (RNN, LSTM)
- Module 8- Transformers (Attention is all you need!)
- Module 9- Prophet and Neural Prophet

