

Comparison of Forecasting Strategies in Machine Learning and Statistical Methods

Author: Dr. Pedram Jahangiry **Course:** DATA5630 - Fall 2025 **Module:** Machine Learning Forecasting

Comparison of Forecasting Strategies in MLForecast (using Random Forest models)

Forecasting Strategy	Forecasting Horizon Handling	Number of Models Trained	Error Propagation Risk	API Support in MLForecast	Feature Construction Notes	Training Complexity	Typical Use Cases
Single-Output (Recursive)	Predicts 1 step ahead and feeds each prediction back in to forecast the next, iterating through the horizon.	1 (single model reused for all steps).	High – errors accumulate at each step as predictions are reused.	Yes – default strategy (no special configuration needed).	Uses standard lag features for the next-step prediction; features are recomputed after each new forecast.	Low – only one model to train (simple pipeline).	Common baseline approach; default in traditional models like ARIMA/ETS; suitable when data is limited or model simplicity is preferred.
Multi-Output (Direct)	Forecast each horizon step with its own independent model (no recursive dependence on prior predictions).	H models (one per forecast horizon step).	Minimal – no accumulation since each step is predicted independently (no feedback of previous forecast errors).	Yes – supported via <code>max_horizon</code> parameter to train one model per step.	Similar lag-based features can be used for each model; no horizon index needed as each model targets a specific lead time.	High – training multiple models is time-consuming (more training time, but can improve accuracy).	Often used when each forecast step's accuracy is critical and extra training is acceptable; ideal for short to mid-range horizons to avoid error compounding.

Forecasting Strategy	Forecasting Horizon Handling	Number of Models Trained	Error Propagation Risk	API Support in MLForecast	Feature Construction Notes	Training Complexity	Typical Use Cases
Multi-output (Single model)	Use a single model to predict all future steps in one shot (multi-output prediction of the entire horizon).	1 (one model outputs a vector of H forecasts).	None – forecasts are generated simultaneously, so no iterative error feedback between steps.	Not directly built-in – requires a custom approach (e.g. add a "horizon" index as a feature or use a multi-target regressor) to use one model for all steps.	May include a horizon indicator feature so the model knows which step is being predicted; no need to recalculate features for each step after training (the model learns all steps together).	Moderate – one complex model to train (can capture dependencies between outputs; may train slower and needs more data to generalize).	Used to avoid error accumulation while managing a single model; useful if ample data is available to learn multi-step patterns. Often seen as a trade-off strategy to achieve direct-accuracy with one model.

Comparison of Forecasting Strategies in StatsForecast

Forecasting Strategy	Description	Supported in StatsForecast (ARIMA/ETS)?	Implementation Details
Single-Output (Recursive)	Train one model for one-step-ahead, then iteratively feed forecasts back in to reach H steps. Avoids retraining, but errors compound.	Yes (default). All statsforecast statistical models use recursive multi-step forecasts by default. For example, calling <code>sf.predict(h=H)</code> on an ARIMA will internally generate 1-step, 2-step, ... forecasts in sequence	Out-of-the-box: simply fit once and use the built-in <code>.predict(h)</code> or <code>.forecast()</code> method to get H steps. The library handles the recursion internally.
Multi-Output (Direct)	Train H separate models, each specialized to predict the t+k horizon directly (one model per step). Eliminates cascading errors at the cost of more models.	No (not built-in). StatsForecast has no parameter to auto-train multiple horizon-specific models – classical ARIMA/ETS are fit on one-step forecasts only. (By contrast, Nixtla's MLForecast supports this via <code>max_horizon</code> for ML models)	Manual only: If desired, you must manually fit separate StatsForecast models or rerun <code>.fit()</code> for each horizon using modified data. This is labor-intensive and not natively supported by the library's API.
Multi-output (Single model)	Train one model that outputs a vector of all H future values in one go. Often implemented via a multi-output regressor or by using horizon as an input feature to a single model.	No. Classical models in statsforecast are univariate and produce a single-step forecast. They cannot natively emit multiple-step outputs in one prediction. There is no ARIMA/ETS that directly yields, say, a 7-day forecast vector – it's achieved through recursion instead	Not supported: There is no out-of-the-box multi-output mode for ARIMA/ETS in StatsForecast. Achieving this would require using a different framework or approach (e.g. ML models with horizon features or specialized multi-output algorithms, which fall outside statsforecast's scope).

Key Takeaways

1. **MLForecast** provides flexible strategies for machine learning models (like Random Forest):

- o Recursive forecasting is the default and simplest approach
- o Direct multi-output forecasting is supported via `max_horizon` parameter
- o Custom single-model multi-output approaches require additional implementation

2. **StatsForecast** focuses on traditional statistical models (ARIMA/ETS):

- o Only recursive forecasting is natively supported
- o Multi-output strategies require manual implementation
- o Best suited for classical time series approaches

3. **Error Propagation** is a critical consideration:

- o Recursive methods compound errors over longer horizons
- o Direct methods avoid error accumulation but require more computational resources
- o Single-model multi-output strikes a balance but needs sufficient training data