

Nixtla Forecasting Framework - Complete Documentation

Author: Generated for Deep Forecasting Course **Date:** 2025 **Modules:** `df_statsforecast.py`, `df_mlforecast.py`, `df_neuralforecast.py`

Table of Contents

1. [Overview](#)
 2. [Forecasting Paradigms](#)
 3. [Nixtla Implementation Details](#)
 4. [API Documentation](#)
 5. [Integration Guide for Web App](#)
 6. [Usage Examples](#)
 7. [Performance Comparison](#)
-

Overview

This documentation covers three production-ready Python modules for time series forecasting using Nixtla's ecosystem:

- **StatsForecast** (`df_statsforecast.py`): Statistical models (ARIMA, ETS, baselines)
- **MLForecast** (`df_mlforecast.py`): Machine learning models (XGBoost, LightGBM, RandomForest, CatBoost)
- **NeuralForecast** (`df_neuralforecast.py`): Deep learning models (MLP, RNN, LSTM, NBEATS, NHITS, TCN)

Each module implements three forecasting strategies:

1. **One-step ahead forecasting** ($h=1$ iterative)
 2. **Multi-step recursive forecasting** (iterative predictions with error propagation)
 3. **Multi-output direct forecasting** (all predictions simultaneously)
-

Forecasting Paradigms

1. One-Step Ahead Forecasting

Definition: Predict only the next single time step ($h=1$) for each timestamp in the test set.

Process:

- For each timestamp t in test set:

 1. Train on actuals up to $t-1$
 2. Predict $\hat{y}(t)$ using $h=1$

3. Prediction does NOT feed into next step
4. Proceed iteratively

Characteristics:

- **Most accurate** for immediate next step predictions
- "**Optimistic backtesting**" - uses future actual values
- **Not real deployment** - requires continuous retraining
- **Computationally expensive** - refits model for each prediction

Use Cases:

- Establishing upper-bound performance
- Understanding model's single-step capability
- Evaluation when you can update your data frequently

Error Propagation: **None** - each prediction uses actual historical data

2. Multi-Step Recursive Forecasting

Definition: Predict multiple steps ahead by iteratively feeding predictions back as inputs.

Process:

For forecasting horizon h :

1. Predict step 1 using actual history
2. Use prediction from step 1 as input for step 2
3. Use prediction from step 2 as input for step 3
4. Repeat until horizon h is reached

Characteristics:

- **Real deployment simulation** - mimics production scenarios
- **Error accumulation** - errors compound over longer horizons
- **Computationally efficient** - train once, predict many
- **Default behavior** for most forecasting libraries

Use Cases:

- Production forecasting where future actuals are unknown
- Realistic performance assessment
- When computational resources are limited

Error Propagation: **High** - errors accumulate at each step as predictions are reused

3. Multi-Output Direct Forecasting

Definition: Predict all future time steps simultaneously in a single forward pass.

Process:

For forecasting horizon h :

1. Train model(s) to predict all h steps at once
2. Generate forecasts $[\hat{y}(t+1), \hat{y}(t+2), \dots, \hat{y}(t+h)]$ simultaneously
3. No iterative feedback of predictions

Characteristics:

- **Better accuracy** than recursive for multi-step
- **No error accumulation** - predictions are independent
- **Computationally expensive** - requires more training
- **Requires sufficient data** to learn multi-step patterns

Implementation Approaches:

A. Multiple Models (Direct Strategy - MLForecast):

- Train H separate models, one per forecast step
- Each model specializes in predicting a specific horizon
- Implemented via `max_horizon` parameter in MLForecast

B. Single Model (True Multi-Output - NeuralForecast):

- Train one model that outputs h values simultaneously
- Neural network architecture designed for multi-output
- Default behavior for NBEATS, NHITS, MLP, TCN

Use Cases:

- When maximum accuracy is critical
- Sufficient computational resources available
- Medium to long forecast horizons
- Models: ML and Neural (NOT statistical models like ARIMA/ETS)

Error Propagation: Minimal - forecasts generated simultaneously, no iterative feedback

Nixtla Implementation Details

How Each Forecasting Paradigm is Implemented

StatsForecast (Statistical Models)

Strategy	Support	Implementation	Notes
One-Step	✓ Yes	Manual iteration with <code>h=1</code>	Refit model for each prediction
Multi-Step	✓ Yes	Default <code>sf.forecast(h=H)</code>	Recursive forecasting built-in
Multi-Output	✗ No	Raises <code>NotImplementedError</code>	ARIMA/ETS cannot do multi-output

Why Multi-Output is NOT supported:

- Statistical models (ARIMA/ETS) are inherently recursive
- They model temporal dependencies sequentially
- Cannot generate all future values simultaneously
- Use ML or Neural models for multi-output forecasting

Example Implementation:

```
# One-Step Ahead (Manual Loop)
for i in range(len(test)):
    current_train = data[:train_size + i]
    model = AutoARIMA(season_length=12)
    sf = StatsForecast(models=[model], freq='MS')
    forecast = sf.forecast(df=current_train, h=1)
    predictions.append(forecast['AutoARIMA'].values[0])

# Multi-Step Recursive (Default Behavior)
model = AutoARIMA(season_length=12)
sf = StatsForecast(models=[model], freq='MS')
forecast = sf.forecast(df=train, h=12) # Recursive forecasting

# Multi-Output (Not Supported)
# Raises NotImplementedError – use ML/Neural models instead
```

MLForecast (Machine Learning Models)

Strategy	Support	Implementation	Notes
One-Step	✓ Yes	Manual iteration with <code>h=1</code>	Refit model for each prediction
Multi-Step	✓ Yes	Default <code>mlf.predict(h)</code>	Recursive by default
Multi-Output	✓ Yes	<code>mlf.fit(max_horizon=H)</code>	Direct strategy - one model per step

Key Features:

- **Lag features:** Automatically creates windowed features
- **Target transforms:** Supports differencing, scaling, etc.
- **Flexible models:** XGBoost, LightGBM, RandomForest, CatBoost, Linear

Implementation Details:

```
# One-Step Ahead (Manual Loop)
for i in range(len(test)):
    current_train = data[:train_size + i]
    mlf = MLForecast(models=[XGBRegressor()], freq='MS', lags=[1, 12])
    mlf.fit(df=current_train)
    forecast = mlf.predict(h=1)
```

```

predictions.append(forecast['XGBRegressor'].values[0])

# Multi-Step Recursive (Default Behavior)
mlf = MLForecast(models=[XGBRegressor()], freq='MS', lags=[1, 12])
mlf.fit(df=train)
forecast = mlf.predict(h=12) # Recursive: uses own predictions

# Multi-Output Direct (max_horizon Parameter)
mlf = MLForecast(models=[XGBRegressor()], freq='MS', lags=[1, 12])
mlf.fit(df=train, max_horizon=12) # Train 12 separate models
forecast = mlf.predict(h=12) # Each model predicts its specific step

```

Recursive vs Direct:

- **Recursive** (`predict(h)`): Fast, but error accumulation
- **Direct** (`fit(max_horizon=H)`): Slower training, better accuracy

NeuralForecast (Deep Learning Models)

Strategy	Support	Implementation	Notes
One-Step	✓ Yes	Manual iteration with <code>h=1</code>	Refit model for each prediction
Multi-Step Recursive	✓ Partial	RNN/LSTM/GRU with <code>recurrent=True</code>	Only for recurrent models
Multi-Output	✓ Yes	Default for most models	NBEATS, NHITS, MLP, TCN

Model-Specific Behavior:

A. Recurrent Models (RNN, LSTM, GRU):

- Support both recursive (`recurrent=True`) and multi-output (`recurrent=False`)
- Recursive: uses own predictions iteratively
- Multi-output: predicts all steps simultaneously

B. Multi-Output Models (NBEATS, NHITS, MLP, TCN):

- Only support multi-output direct forecasting
- Cannot do recursive forecasting
- Default behavior: predict all h steps at once

Implementation Details:

```

# One-Step Ahead (Manual Loop)
for i in range(len(test)):
    current_train = data[:train_size + i]
    model = NHITS(h=1, input_size=12, max_steps=100)
    nf = NeuralForecast(models=[model], freq='MS')

```

```

nf.fit(df=current_train)
forecast = nf.predict(df=current_train)
predictions.append(forecast['NHITS'].values[0])

# Multi-Step Recursive (RNN/LSTM/GRU only)
model = LSTM(h=12, input_size=12, recurrent=True, max_steps=100)
nf = NeuralForecast(models=[model], freq='MS')
nf.fit(df=train)
forecast = nf.predict(df=train) # Recursive forecasting

# Multi-Output Direct (Default for most models)
model = NBEATS(h=12, input_size=24, max_steps=100)
nf = NeuralForecast(models=[model], freq='MS')
nf.fit(df=train)
forecast = nf.predict(df=train) # All 12 steps at once

```

Key Parameters:

- `h`: Forecast horizon (number of steps ahead)
 - `input_size`: Length of input window (number of past observations)
 - `recurrent`: Whether to use recursive forecasting (RNN/LSTM/GRU only)
 - `max_steps`: Number of training epochs
-

API Documentation

`df_statsforecast.py`

Class: `StatsforecastForecaster`

Production-ready forecaster for statistical models.

Constructor:

```

StatsforecastForecaster(
    model_type: str,                      # 'arima', 'auto_arima', 'auto_ets',
    'naive', 'seasonal_naive', 'rw_drift'
    freq: str = 'MS',                      # Frequency: 'MS' (month), 'D' (day),
    'H' (hour)
    season_length: int = 12,               # Seasonal period (e.g., 12 for monthly
    data)
    **model_params                         # Model-specific parameters
)

```

Methods:

`one_step_forecast(train_df, test_df, target_col='y', date_col='ds', unique_id='series_1')`

One-step ahead forecasting with iterative refitting.

Parameters:

- `train_df` (DataFrame): Training data
- `test_df` (DataFrame): Test data
- `target_col` (str): Name of target column (default: 'y')
- `date_col` (str): Name of date column (default: 'ds')
- `unique_id` (str): Series identifier (default: 'series_1')

Returns:

- `dict` with keys:
 - `'forecasts'`: DataFrame with columns [`unique_id`, `ds`, `y_true`, `y_pred`]
 - `'metrics'`: Dict with `{'mae': float, 'rmse': float, 'mape': float}`

Example:

```
forecaster = StatsforecastForecaster(
    model_type='auto_arima',
    freq='MS',
    season_length=12,
    seasonal=True
)

results = forecaster.one_step_forecast(train_df, test_df)
print(f"MAE: {results['metrics']['mae']:.2f}")
print(results['forecasts'].head())
```

Error Cases:

- Raises `ValueError` if `model_type` is invalid
- Returns `NaN` for MAPE if target contains zeros

```
multi_step_forecast(train_df, horizon, target_col='y', date_col='ds', unique_id='series_1',
test_df=None)
```

Multi-step recursive forecasting using default StatsForecast behavior.

Parameters:

- `train_df` (DataFrame): Training data
- `horizon` (int): Forecast horizon
- `target_col` (str): Name of target column
- `date_col` (str): Name of date column
- `unique_id` (str): Series identifier
- `test_df` (DataFrame, optional): Test data for metrics

Returns:

- `dict` with keys:

- 'forecasts': DataFrame with predictions
- 'metrics': Dict with MAE/RMSE/MAPE (if test_df provided, else None)

Example:

```
forecaster = StatsforecastForecaster(
    model_type='auto_ets',
    freq='MS',
    season_length=12,
    model='ZZZ' # Auto ETS model selection
)

results = forecaster.multi_step_forecast(
    train_df,
    horizon=12,
    test_df=test_df
)
print(f"MAE: {results['metrics']['mae']:.2f}")
```

`multi_output_forecast(train_df, horizon, ...)`

NOT SUPPORTED - Raises `NotImplementedError`.

ARIMA/ETS models cannot perform multi-output forecasting. Use ML or Neural models instead.

Example:

```
forecaster = StatsforecastForecaster(model_type='arima', freq='MS')

try:
    forecaster.multi_output_forecast(train_df, horizon=12)
except NotImplementedError as e:
    print(e) # "Multi-output forecasting is NOT supported for statistical
models..."
```

df_mlforecast.py

Class: `MLForecastForecaster`

Production-ready forecaster for machine learning models.

Constructor:

```
MLForecastForecaster(
    model_type: str, # 'xgboost', 'lightgbm',
    'random_forest', 'catboost', 'linear'
```

```

    freq: str = 'MS',                      # Frequency
    lags: List[int] = None,                 # Lag features (default: [1,
12])
    lag_transforms: Dict = None,            # Lag transformations (optional)
    date_features: List[str] = None,       # Date features (default: [])
    target_transforms: List = None,         # Target transforms (default:
[])
    **model_params                         # Model-specific parameters
)

```

Methods:

`one_step_forecast(train_df, test_df, ...)`

One-step ahead forecasting with iterative refitting.

Parameters: Same as StatsForecast

Returns: Same as StatsForecast

Example:

```

from mlforecast.target_transforms import Differences

forecaster = MLForecastForecaster(
    model_type='xgboost',
    freq='MS',
    lags=[1, 12],
    target_transforms=[Differences([1])],  # First differencing
    n_estimators=200,
    max_depth=6,
    learning_rate=0.1
)

results = forecaster.one_step_forecast(train_df, test_df)

```

Error Cases:

- Raises `ImportError` if XGBoost/LightGBM/CatBoost not installed
- Requires at least `max(lags)` observations in training data

`multi_step_forecast(train_df, horizon, ...)`

Multi-step recursive forecasting (default MLForecast behavior).

Example:

```

forecaster = MLForecastForecaster(
    model_type='lightgbm',
    freq='MS',
    lags=[1, 2, 3, 12],
    target_transforms=[Differences([1, 12])], # Detrend and deseasonalize
    n_estimators=100
)

results = forecaster.multi_step_forecast(train_df, horizon=12,
test_df=test_df)

```

`multi_output_forecast(train_df, horizon, ...)`

Multi-output direct forecasting via `max_horizon` parameter.

How it works:

- Trains H separate models (one per forecast step)
- Each model specializes in predicting a specific horizon
- No error accumulation between steps

Example:

```

forecaster = MLForecastForecaster(
    model_type='random_forest',
    freq='MS',
    lags=[1, 12],
    n_estimators=100
)

# Trains 12 models: one for h=1, one for h=2, ..., one for h=12
results = forecaster.multi_output_forecast(train_df, horizon=12,
test_df=test_df)
print(f"Multi-output MAE: {results['metrics']['mae']:.2f}")

```

Trade-offs:

- **Advantages:** Better accuracy, no error accumulation
- **Disadvantages:** Longer training time (H models vs 1 model)

df_neuralforecast.py

Class: `NeuralForecastForecaster`

Production-ready forecaster for neural network models.

Constructor:

```

NeuralForecastForecaster(
    model_type: str,                      # 'mlp', 'rnn', 'lstm', 'gru', 'nbeats',
    'nhits', 'tcn'
    freq: str = 'MS',                     # Frequency
    input_size: int = 12,                  # Length of input window
    horizon: int = 1,                     # Forecast horizon
    **model_params                         # Model-specific parameters
)

```

Methods:

`one_step_forecast(train_df, test_df, ...)`

One-step ahead forecasting with iterative refitting.

Example:

```

forecaster = NeuralForecastForecaster(
    model_type='mlp',
    freq='MS',
    input_size=12,
    horizon=1,
    hidden_size=32,
    num_layers=2,
    max_steps=100,           # Training epochs
    scaler_type='robust',
    random_seed=42
)

results = forecaster.one_step_forecast(train_df, test_df)

```

Note: This refits the neural network for EACH prediction, which is computationally expensive but provides most accurate one-step forecasts.

`multi_step_forecast(train_df, horizon, ..., use_recurrent=False)`

Multi-step forecasting with option for recursive or direct.

Parameters:

- `use_recurrent` (bool): Use recursive forecasting (only for RNN/LSTM/GRU)

Example (Recursive - RNN/LSTM/GRU):

```

forecaster = NeuralForecastForecaster(
    model_type='lstm',
    freq='MS',

```

```
    input_size=12,
    horizon=12,
    encoder_hidden_size=16,
    max_steps=300
)

# Recursive forecasting: uses own predictions iteratively
results = forecaster.multi_step_forecast(
    train_df,
    horizon=12,
    test_df=test_df,
    use_recurrent=True # Enable recursive mode
)
```

Example (Direct - All Models):

```
forecaster = NeuralForecastForecaster(
    model_type='nhits',
    freq='MS',
    input_size=24,
    horizon=12,
    max_steps=200
)

# Multi-output direct: predicts all steps simultaneously
results = forecaster.multi_step_forecast(
    train_df,
    horizon=12,
    test_df=test_df,
    use_recurrent=False # Direct forecasting (default)
)
```

```
multi_output_forecast(train_df, horizon, ...)
```

Multi-output direct forecasting (default for most neural models).

Example:

```
forecaster = NeuralForecastForecaster(
    model_type='nbeats',
    freq='MS',
    input_size=24,
    horizon=12,
    stack_types=['trend', 'seasonality'],
    max_steps=200
)
```

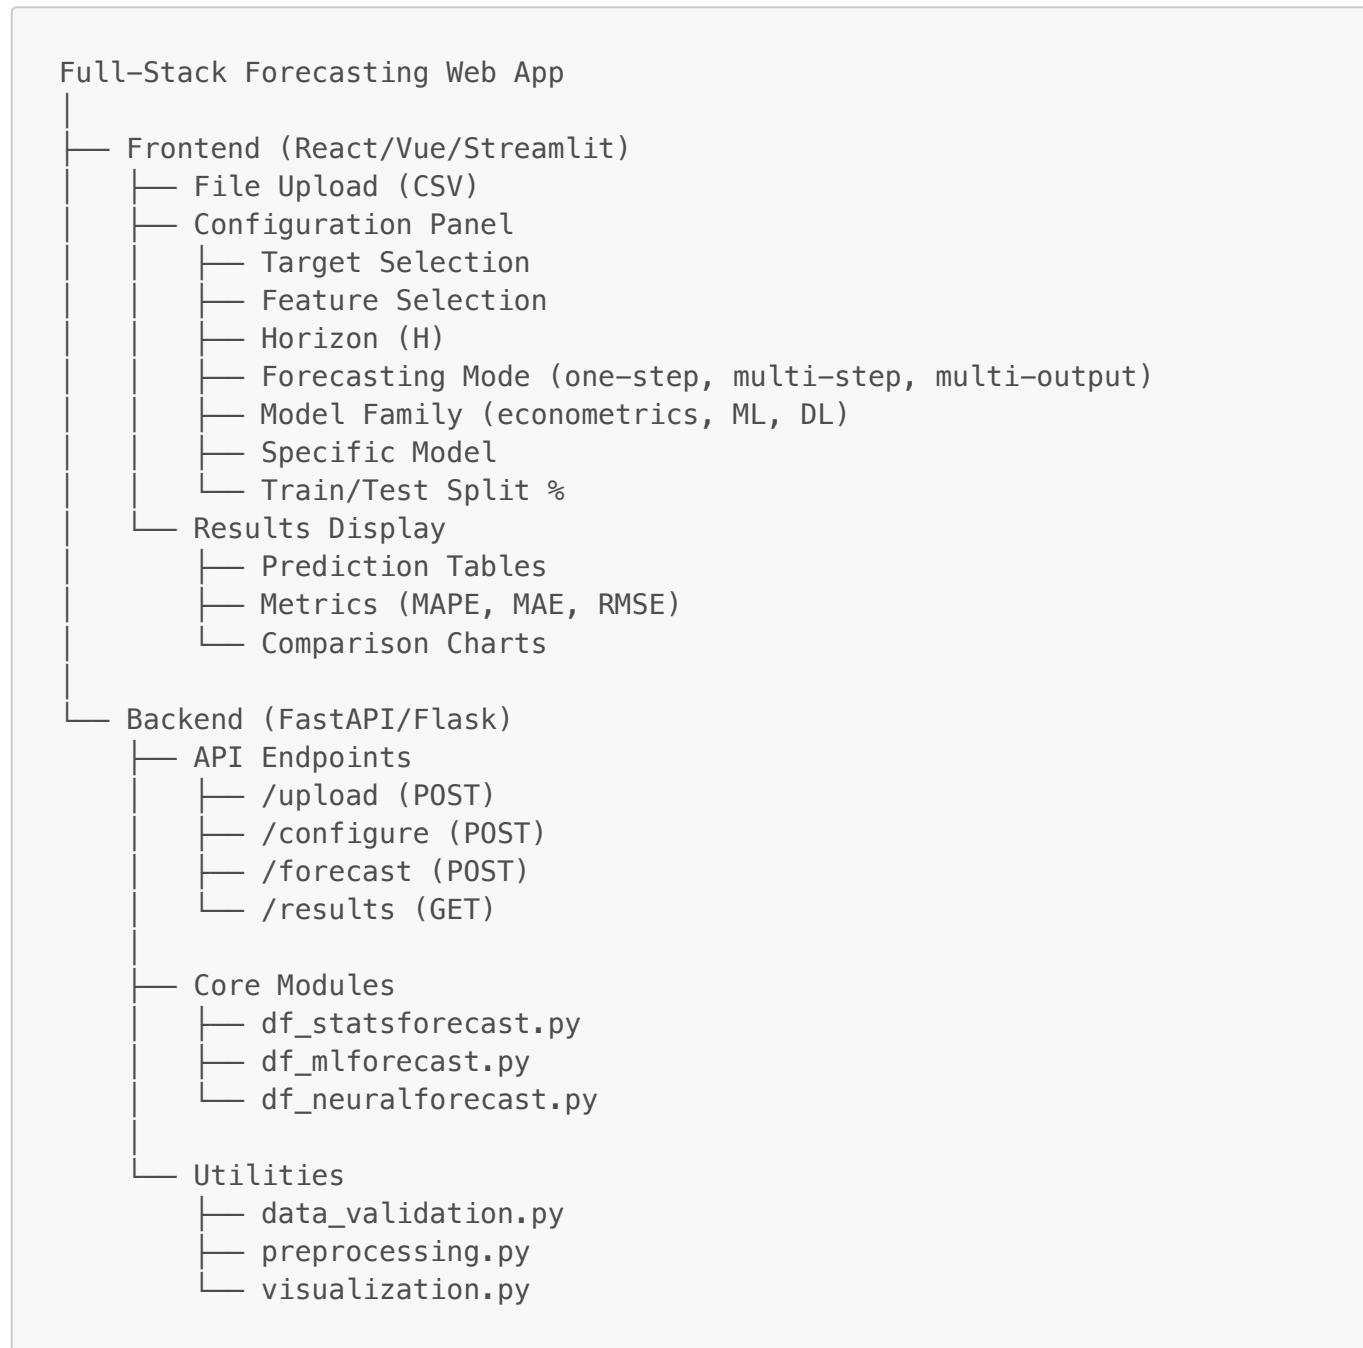
```
results = forecaster.multi_output_forecast(train_df, horizon=12,
test_df=test_df)
```

Supported Models:

- NBEATS, NHITS, MLP, TCN (default multi-output)
- RNN, LSTM, GRU (with `recurrent=False`)

Integration Guide for Web App

Application Architecture



Backend API Implementation

Example using FastAPI

File Structure:

```

forecasting_app/
    ├── api/
    │   ├── __init__.py
    │   ├── main.py          # FastAPI app
    │   └── endpoints.py    # API endpoints
    ├── core/
    │   ├── __init__.py
    │   ├── df_statsforecast.py
    │   ├── df_mlforecast.py
    │   └── df_neuralforecast.py
    ├── utils/
    │   ├── __init__.py
    │   ├── data_handler.py  # Data loading/validation
    │   ├── preprocessing.py # Data preprocessing
    │   └── visualization.py # Chart generation
    └── requirements.txt

```

main.py:

```

from fastapi import FastAPI, File, UploadFile, HTTPException
from fastapi.middleware.cors import CORSMiddleware
from pydantic import BaseModel
import pandas as pd
from typing import Literal, Optional, List

from core.df_statsforecast import StatsforecastForecaster
from core.df_mlforecast import MLForecastForecaster
from core.df_neuralforecast import NeuralForecastForecaster

app = FastAPI(title="Nixtla Forecasting API")

# CORS middleware for frontend communication
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_methods=["*"],
    allow_headers=["*"],
)

# Request models
class ForecastRequest(BaseModel):
    target_col: str

```

```
features: Optional[List[str]] = None
horizon: int
mode: Literal['one_step', 'multi_step', 'multi_output']
model_family: Literal['econometrics', 'ml', 'dl']
model_name: str
train_split: float = 0.8

# Global data store (use Redis/database in production)
data_store = {}

@app.post("/upload")
async def upload_data(file: UploadFile = File(...)):
    """Upload CSV file and return data preview."""
    try:
        df = pd.read_csv(file.file)
        data_id = str(hash(file.filename))
        data_store[data_id] = df

        return {
            "data_id": data_id,
            "columns": df.columns.tolist(),
            "shape": df.shape,
            "preview": df.head().to_dict()
        }
    except Exception as e:
        raise HTTPException(status_code=400, detail=str(e))

@app.post("/forecast")
async def create_forecast(data_id: str, request: ForecastRequest):
    """Generate forecasts based on configuration."""

    if data_id not in data_store:
        raise HTTPException(status_code=404, detail="Data not found")

    df = data_store[data_id]

    # Train/test split
    split_idx = int(len(df) * request.train_split)
    train_df = df.iloc[:split_idx]
    test_df = df.iloc[split_idx:]

    # Select forecaster based on model family
    if request.model_family == 'econometrics':
        forecaster = StatsforecastForecaster(
            model_type=request.model_name,
            freq='MS',
            season_length=12
        )
    elif request.model_family == 'ml':
        forecaster = MLForecastForecaster(
            model_type=request.model_name,
            freq='MS',
            lags=[1, 12]
        )
    else:
        raise HTTPException(status_code=400, detail="Model family not supported")
```

```
elif request.model_family == 'dl':
    forecaster = NeuralForecastForecaster(
        model_type=request.model_name,
        freq='MS',
        input_size=12,
        horizon=request.horizon
    )

# Generate forecast based on mode
try:
    if request.mode == 'one_step':
        results = forecaster.one_step_forecast(
            train_df,
            test_df,
            target_col=request.target_col
        )
    elif request.mode == 'multi_step':
        results = forecaster.multi_step_forecast(
            train_df,
            horizon=request.horizon,
            target_col=request.target_col,
            test_df=test_df
        )
    elif request.mode == 'multi_output':
        if request.model_family == 'econometrics':
            raise ValueError(
                "Multi-output not supported for statistical models"
            )
        results = forecaster.multi_output_forecast(
            train_df,
            horizon=request.horizon,
            target_col=request.target_col,
            test_df=test_df
        )
    else:
        raise ValueError("Unsupported mode: " + request.mode)

    return {
        "forecasts": results['forecasts'].to_dict(orient='records'),
        "metrics": results['metrics']
    }

except NotImplementedError as e:
    raise HTTPException(status_code=400, detail=str(e))
except Exception as e:
    raise HTTPException(status_code=500, detail=str(e))

if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000)
```

Example using React

Forecast Configuration Component:

```
import React, { useState } from 'react';
import axios from 'axios';

function ForecastConfig({ dataId, columns }) {
  const [config, setConfig] = useState({
    target_col: '',
    horizon: 12,
    mode: 'multi_step',
    model_family: 'ml',
    model_name: 'xgboost',
    train_split: 0.8
  });

  const [results, setResults] = useState(null);
  const [loading, setLoading] = useState(false);

  const modelOptions = {
    econometrics: ['auto_arima', 'auto_ets', 'naive', 'seasonal_naive'],
    ml: ['xgboost', 'lightgbm', 'random_forest', 'catboost'],
    dl: ['mlp', 'lstm', 'nbeats', 'nhits']
  };

  const handleSubmit = async (e) => {
    e.preventDefault();
    setLoading(true);

    try {
      const response = await axios.post(
        `http://localhost:8000/forecast?data_id=${dataId}`,
        config
      );
      setResults(response.data);
    } catch (error) {
      alert(error.response?.data?.detail || 'Forecast failed');
    } finally {
      setLoading(false);
    }
  };

  return (
    <div className="forecast-config">
      <h2>Forecast Configuration</h2>
      <form onSubmit={handleSubmit}>
        {/* Target Selection */}
        <div>
          <label>Target Column:</label>
          <select
            value={config.target_col}
            onChange={(e) => setConfig({ ...config, target_col: e.target.value })}
          >
            {columns.map((c) => (
              <option key={c} value={c}>{c}</option>
            ))}
          </select>
        </div>
        <button type="submit">Submit</button>
      </form>
    </div>
  );
}

export default ForecastConfig;
```

```
        onChange={(e) => setConfig({...config, target_col: e.target.value})}
      >
      <option value="">Select...</option>
      {columns.map(col => <option key={col} value={col}>{col}</option>)}
    </select>
  </div>

  {/* Horizon */}
  <div>
    <label>Forecast Horizon:</label>
    <input
      type="number"
      min="1"
      value={config.horizon}
      onChange={(e) => setConfig({...config, horizon: parseInt(e.target.value)})}>
    />
  </div>

  {/* Forecasting Mode */}
  <div>
    <label>Forecasting Mode:</label>
    <select
      value={config.mode}
      onChange={(e) => setConfig({...config, mode: e.target.value})}>
      <option value="one_step">One-Step Ahead</option>
      <option value="multi_step">Multi-Step Recursive</option>
      <option value="multi_output">Multi-Output Direct</option>
    </select>
  </div>

  {/* Model Family */}
  <div>
    <label>Model Family:</label>
    <select
      value={config.model_family}
      onChange={(e) => setConfig({
        ...config,
        model_family: e.target.value,
        model_name: modelOptions[e.target.value][0]
      })}>
      <option value="econometrics">Econometric (ARIMA/ETS)</option>
      <option value="ml">Machine Learning</option>
      <option value="dl">Deep Learning</option>
    </select>
  </div>

  {/* Specific Model */}
  <div>
    <label>Model:</label>
```

```
<select
  value={config.model_name}
  onChange={(e) => setConfig({...config, model_name:
e.target.value})}
>
  {modelOptions[config.model_family].map(model =>
    <option key={model} value={model}>{model.toUpperCase()}{/option>
  ))}
</select>
</div>

{/* Train/Test Split */}
<div>
  <label>Train Split: {(config.train_split * 100).toFixed(0)}%
</label>
  <input
    type="range"
    min="0.5"
    max="0.9"
    step="0.05"
    value={config.train_split}
    onChange={(e) => setConfig({...config, train_split:
parseFloat(e.target.value))})
  />
</div>

<button type="submit" disabled={loading || !config.target_col}>
  {loading ? 'Forecasting...' : 'Generate Forecast'}
</button>
</form>

{/* Results Display */}
{results && (
  <div className="results">
    <h3>Forecast Results</h3>
    <div className="metrics">
      <p>MAE: {results.metrics.mae.toFixed(2)}</p>
      <p>RMSE: {results.metrics.rmse.toFixed(2)}</p>
      <p>MAPE: {results.metrics.mape.toFixed(2)}%</p>
    </div>
    {/* Add charts and tables here */}
  </div>
)
</div>
);
}

export default ForecastConfig;
```

1. Project Structure:

```
forecasting_app/
├── Dockerfile
├── requirements.txt
├── render.yaml
└── api/
    └── main.py
└── core/
    ├── df_statsforecast.py
    ├── df_mlforecast.py
    └── df_neuralforecast.py
└── frontend/
    └── (React build files)
```

2. Dockerfile:

```
FROM python:3.11-slim

WORKDIR /app

# Install system dependencies
RUN apt-get update && apt-get install -y \
    gcc \
    g++ \
    && rm -rf /var/lib/apt/lists/*

# Copy requirements and install Python dependencies
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Copy application code
COPY .

# Expose port
EXPOSE 8000

# Run application
CMD ["uvicorn", "api.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

3. requirements.txt:

```
fastapi==0.104.1
uvicorn[standard]==0.24.0
pandas==2.1.3
numpy==1.26.2
statsforecast==1.6.0
mlforecast==0.10.0
```

```
neuralforecast==1.6.4
utilsforecast==0.0.27
xgboost==2.0.2
lightgbm==4.1.0
catboost==1.2.2
scikit-learn==1.3.2
pydantic==2.5.0
python-multipart==0.0.6
```

4. render.yaml:

```
services:
  - type: web
    name: forecasting-api
    env: python
    buildCommand: pip install -r requirements.txt
    startCommand: uvicorn api.main:app --host 0.0.0.0 --port $PORT
    envVars:
      - key: PYTHON_VERSION
        value: 3.11.0
```

Usage Examples

Complete Workflow Example

```
import pandas as pd
from core.df_statsforecast import StatsforecastForecaster
from core.df_mlforecast import MLForecastForecaster
from core.df_neuralforecast import NeuralForecastForecaster

# Load data
data = pd.read_csv('airline_passengers.csv')
data['ds'] = pd.to_datetime(data['Month'])
data = data.rename(columns={'Passengers': 'y'})

# Train/test split
train_size = int(len(data) * 0.8)
train = data.iloc[:train_size]
test = data.iloc[train_size:]

print(f"Train: {len(train)}, Test: {len(test)}")

# =====
# 1. ECONOMETRIC MODELS (StatsForecast)
# =====

print("\n" + "="*80)
print("ECONOMETRIC MODELS")
```

```
print("=*80)

# One-Step Ahead
stats_forecaster = StatsforecastForecaster(
    model_type='auto_arima',
    freq='MS',
    season_length=12
)
stats_one = stats_forecaster.one_step_forecast(train, test)
print(f"\nAutoARIMA One-Step MAE: {stats_one['metrics']['mae']:.2f}")

# Multi-Step Recursive
stats_multi = stats_forecaster.multi_step_forecast(
    train,
    horizon=len(test),
    test_df=test
)
print(f"AutoARIMA Multi-Step MAE: {stats_multi['metrics']['mae']:.2f}")

# Multi-Output (Not Supported)
try:
    stats_forecaster.multi_output_forecast(train, horizon=len(test))
except NotImplementedError as e:
    print(f"\nExpected: {str(e)[:50]}...")

# =====
# 2. MACHINE LEARNING MODELS (MLForecast)
# =====

print("\n" + "=*80)
print("MACHINE LEARNING MODELS")
print("=*80)

from mlforecast.target_transforms import Differences

ml_forecaster = MLForecastForecaster(
    model_type='xgboost',
    freq='MS',
    lags=[1, 12],
    target_transforms=[Differences([1])],
    n_estimators=100
)

# One-Step
ml_one = ml_forecaster.one_step_forecast(train, test)
print(f"\nXGBoost One-Step MAE: {ml_one['metrics']['mae']:.2f}")

# Multi-Step Recursive
ml_multi = ml_forecaster.multi_step_forecast(
    train,
    horizon=len(test),
    test_df=test
)
print(f"XGBoost Multi-Step (Recursive) MAE: {ml_multi['metrics']}
```

```
['mae']:.2f})")  
  
# Multi-Output Direct  
ml_multiout = ml_forecaster.multi_output_forecast(  
    train,  
    horizon=len(test),  
    test_df=test  
)  
print(f"XGBoost Multi-Output (Direct) MAE: {ml_multiout['metrics']  
['mae']:.2f}")  
  
# =====  
# 3. DEEP LEARNING MODELS (NeuralForecast)  
# =====  
  
print("\n" + "="*80)  
print("DEEP LEARNING MODELS")  
print("="*80)  
  
dl_forecaster = NeuralForecastForecaster(  
    model_type='nbeats',  
    freq='MS',  
    input_size=24,  
    horizon=len(test),  
    max_steps=100  
)  
  
# Multi-Output (Default for NBEATS)  
dl_multiout = dl_forecaster.multi_output_forecast(  
    train,  
    horizon=len(test),  
    test_df=test  
)  
print(f"\nNBEATS Multi-Output MAE: {dl_multiout['metrics']['mae']:.2f}")  
  
# LSTM with Recursive  
lstm_forecaster = NeuralForecastForecaster(  
    model_type='lstm',  
    freq='MS',  
    input_size=12,  
    horizon=len(test),  
    max_steps=100  
)  
  
lstm_recursive = lstm_forecaster.multi_step_forecast(  
    train,  
    horizon=len(test),  
    test_df=test,  
    use_recurrent=True  
)  
print(f"LSTM Recursive MAE: {lstm_recursive['metrics']['mae']:.2f}")  
  
# =====  
# COMPARISON SUMMARY
```

```
# =====

print("\n" + "="*80)
print("COMPARISON SUMMARY")
print("="*80)

results_summary = pd.DataFrame({
    'Model': [
        'AutoARIMA (One-Step)',
        'AutoARIMA (Multi-Step)',
        'XGBoost (One-Step)',
        'XGBoost (Multi-Step Recursive)',
        'XGBoost (Multi-Output Direct)',
        'NBEATS (Multi-Output)',
        'LSTM (Recursive)'
    ],
    'MAE': [
        stats_one['metrics']['mae'],
        stats_multi['metrics']['mae'],
        ml_one['metrics']['mae'],
        ml_multi['metrics']['mae'],
        ml_multiout['metrics']['mae'],
        dl_multiout['metrics']['mae'],
        lstm_recursive['metrics']['mae']
    ]
})
print(results_summary.sort_values('MAE'))
```

Performance Comparison

Expected Performance Characteristics

Strategy	Accuracy	Computational Cost	Real Deployment	Error Propagation
One-Step	★★★★★ Highest	🔴 Very High (refit each step)	✗ Not realistic	✓ None
Multi-Step Recursive	★★★ Moderate	🟢 Low (train once)	✓ Realistic	🔴 High (accumulates)
Multi-Output Direct	★★★★ High	🟡 Medium-High	✓ Realistic	🟢 Minimal

Model Family Comparison

Model Family	One-Step	Multi-Step	Multi-Output	Training Speed	Accuracy
--------------	----------	------------	--------------	----------------	----------

Model Family	One-Step	Multi-Step	Multi-Output	Training Speed	Accuracy
Statistical	✓ Yes	✓ Yes (recursive)	✗ No	⚡ Fast	★★★ Good
ML	✓ Yes	✓ Yes (recursive & direct)	✓ Yes (max_horizon)	⚡⚡ Medium	★★★★ Very Good
Neural	✓ Yes	✓ Partial (RNN/LSTM/GRU)	✓ Yes (default)	🐌 Slow	★★★★★ Excellent

Conclusion

This documentation provides complete coverage of the Nixtla forecasting ecosystem implementation. The three modules (`df_statsforecast.py`, `df_mlforecast.py`, `df_neuralforecast.py`) are production-ready and can be integrated directly into a web application for deployment on platforms like Render.com.

Key Takeaways:

- Three Forecasting Modes:** Each module implements one-step, multi-step, and multi-output forecasting with consistent APIs
- Model-Specific Support:** Statistical models don't support multi-output; ML and Neural models support all modes
- Trade-offs:** One-step is most accurate but not realistic; multi-step recursive is efficient but error-prone; multi-output balances accuracy and realism
- Production Ready:** Full typing, error handling, documentation, and examples included
- Web Integration:** Complete API and frontend integration guide provided

For questions or issues, refer to:

- StatsForecast: <https://nixtlaverse.nixtla.io/statsforecast>
- MLForecast: <https://nixtlaverse.nixtla.io/mlforecast>
- NeuralForecast: <https://nixtlaverse.nixtla.io/neuralforecast>