Road map!

- Module 1- Introduction to Deep Forecasting
- Module 2- Setting up Deep Forecasting Environment
- Module 3- Exponential Smoothing
- Module 4- ARIMA models
- Module 5- Machine Learning for Time series Forecasting
- Module 6- Deep Neural Networks
- Module 7- Deep Sequence Modeling (RNN, LSTM)
- Module 8- Transformers (Attention is all you need!)
- Module 9- Prophet and Neural Prophet

Module 4 - Part I
Deep Neural Networks
Basics

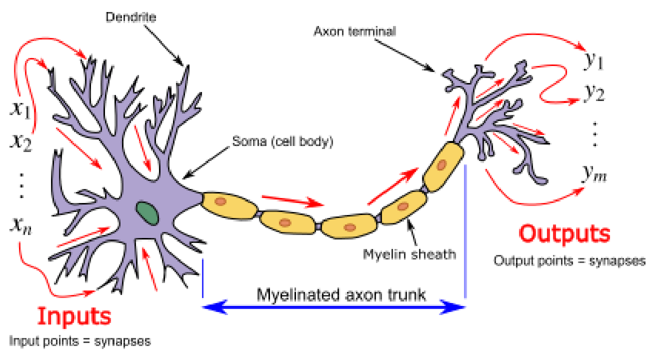

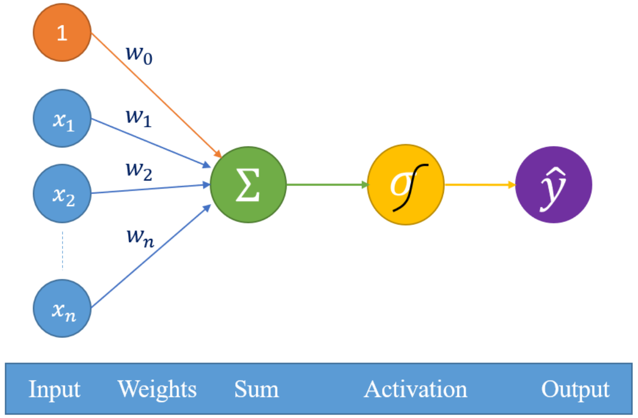

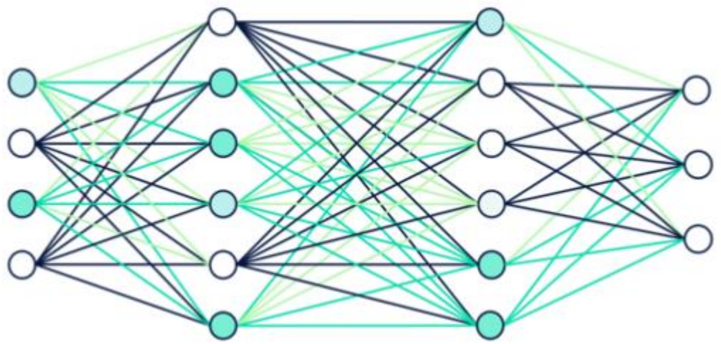

A simple example!

- Consider the housing price prediction! $price = f(sqft)$

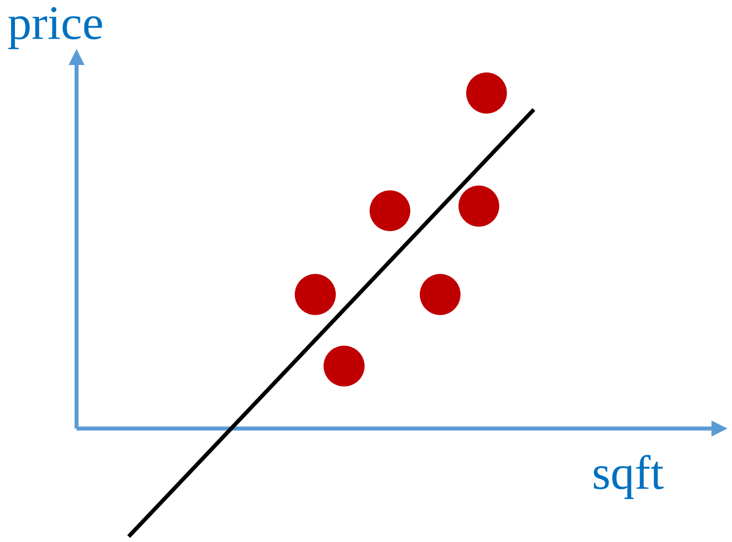

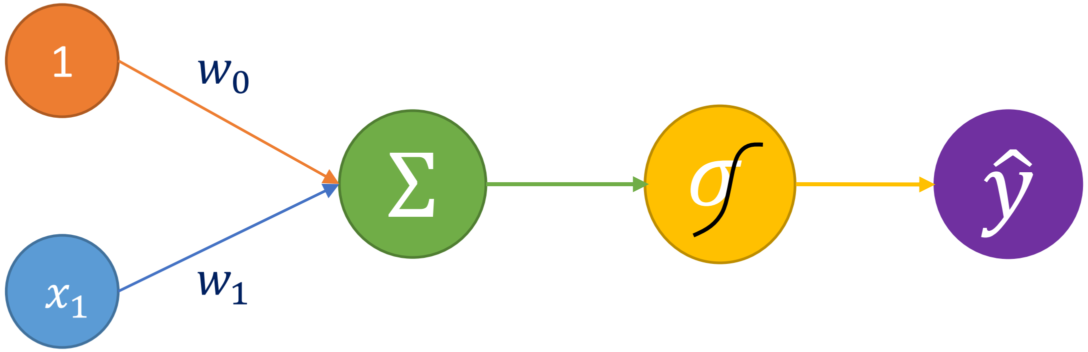

Forward Propagation

- Forward propagation is the process of calculating the output of a neural network, given an input.
- w_0 is the bias term which allows shifting Σ to the left or right.

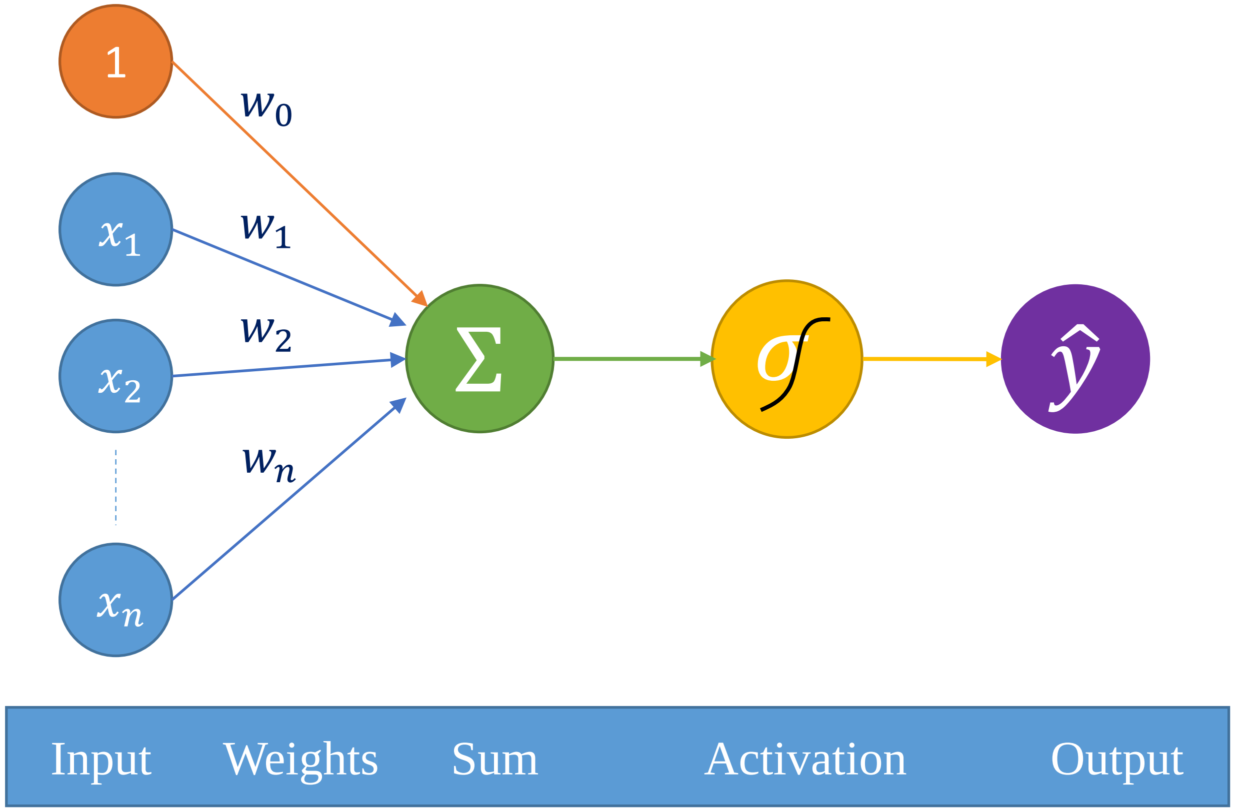

What are Neurons?

- Neurons, are the simplest elements or building blocks in a neural network. They are inspired by biological neurons that are found in the human brain.

- What happens inside the neurons?

Activation function

Bias

$$\hat{y} = \sigma\left(w_0 + \sum_{i=1}^{n} w_i x_i\right)$$

$$\hat{y} = \sigma(w_0 + W^T X)$$

$$\hat{y} = \sigma(Z)$$

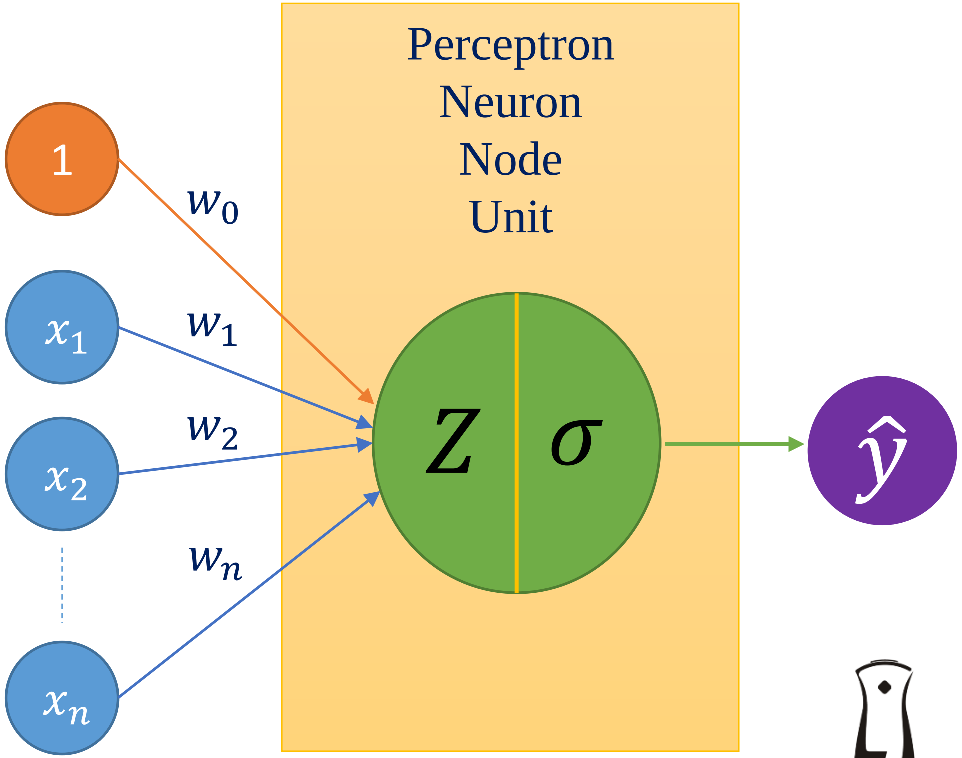

Jon M. Huntsman School of Business
UtahStateUniversity

Pedram Jahangiry

Activation function

- An activation function is a mathematical function that is applied to the output of each neuron in a neural network
- Activation functions allow the network to learn non-linearities and complex patterns from the data and make accurate predictions

Sigmoid	Tanh	ReLU	Leaky ReLU
$g(z) = \frac{1}{1+e^{-z}}$	$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$g(z) = \max(0, z)$	$g(z) = \max(\epsilon z, z)$ with $\epsilon \ll 1$

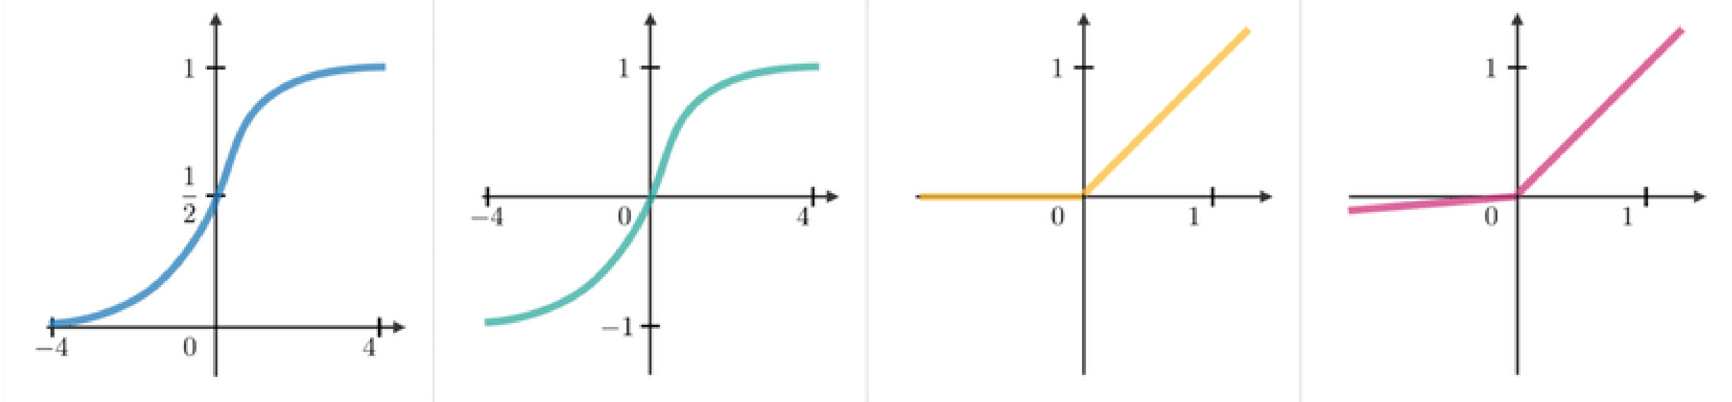

JON M. HUNTSMAN SCHOOL OF BUSINESS
UtahStateUniversity

Pedram Jahangiry

Softmax Activation function

- Softmax activation function is often used in classification tasks, where the goal is to predict which of a fixed set of classes (more than two classes) a particular sample belongs to.
- The Softmax function takes in a vector of real numbers and converts it into a probability distribution, where the sum of all the probabilities is equal to 1.

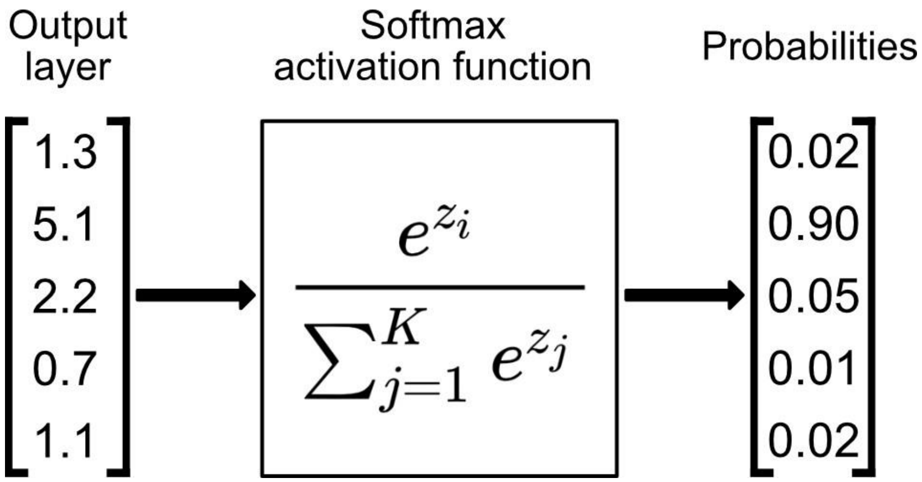

Which activation function?

Again, there is no “right” answer! However:

- For the hidden layers, almost always tanh is better than sigmoid because it center the data for the next layer.
- One downside of both sigmoid or tanh is that the gradient is very small for extreme values.

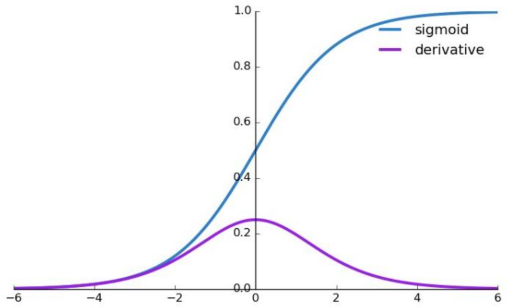

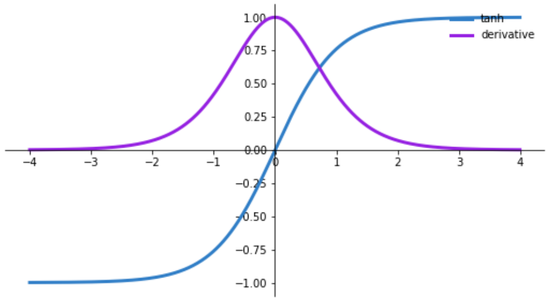

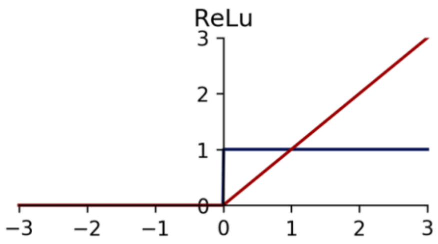

- In practice, RELU works better/faster than sigmoid or tanh for hidden layers.
- Leaky RELU might work better than RELU, but if we have enough number of hidden layers, RELU is just fine.
- Sigmoid is mostly used for output layer!

Jon M. Huntsman School of Business
Utah State University
Pedram Jahangiry

Sigmoid activation function! example

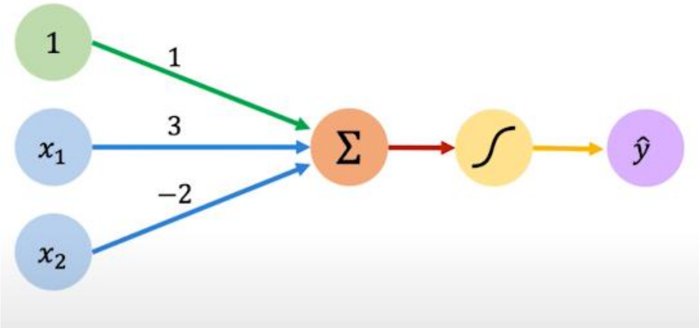

We have: $w_0 = 1$ and $\boldsymbol{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

$$\hat{y} = g\left(w_0 + \boldsymbol{X}^T \boldsymbol{W} \right)$$
$$= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right)$$
$$\hat{y} = g\big(\underbrace{1 + 3x_1 - 2x_2}_{\text{This is just a line in 2D!}}\big)$$

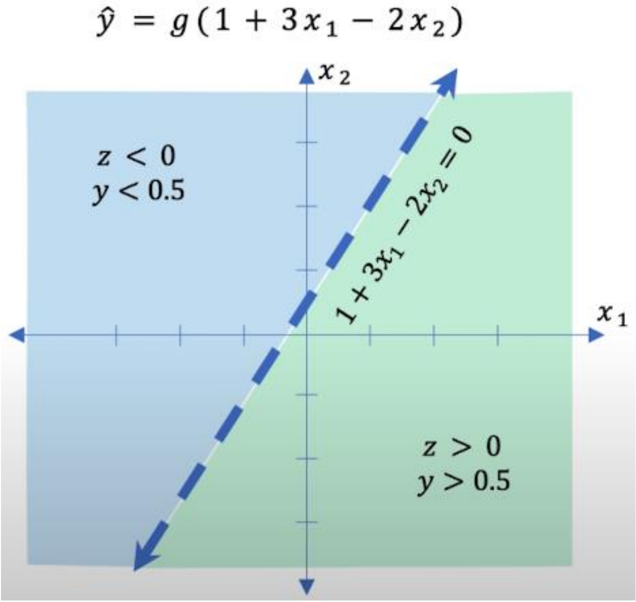

Why do we need non-linear activation functions?

- If we use linear activation function $\sigma(z) = z = w_0 + W^T X$, then the NN will always output a linear function of the inputs regardless of the number of neurons or hidden layers used.
- Linear activation functions produce linear decision boundaries!
- Activation functions allow the network to approximate complex patterns!

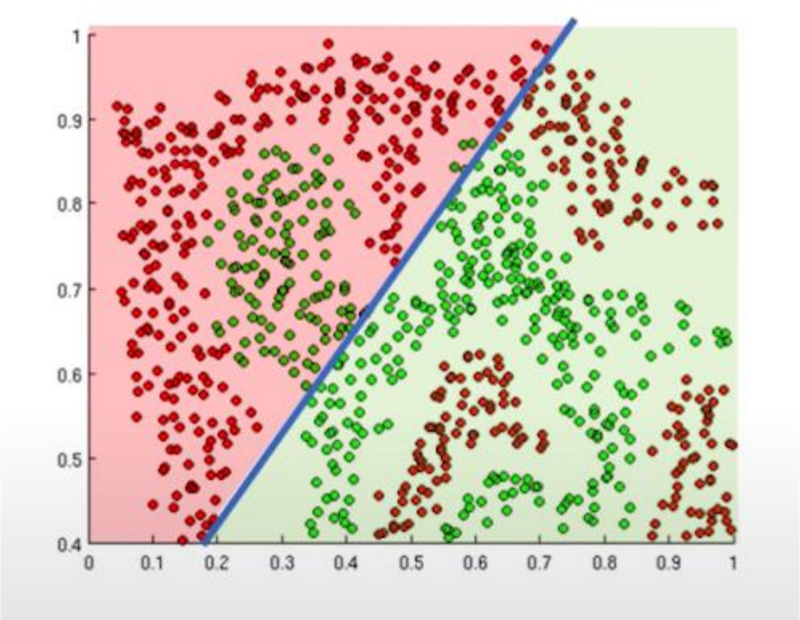

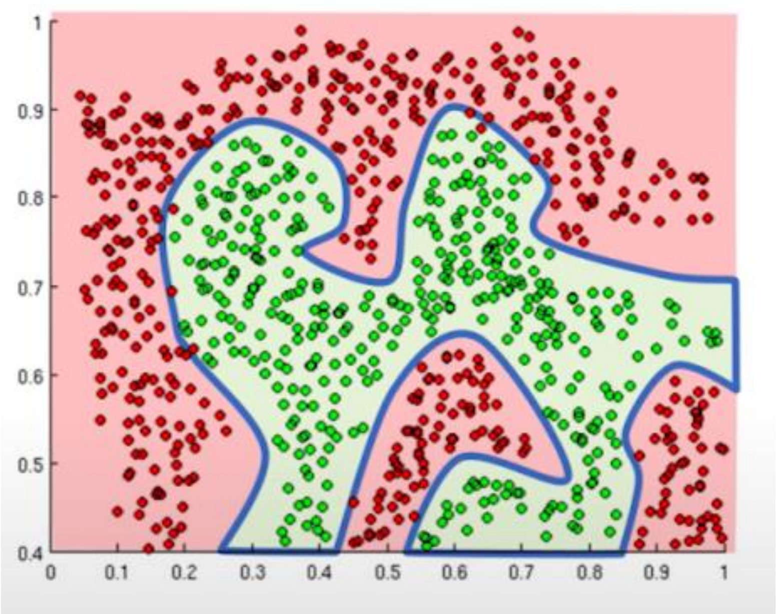

Building a Neural Network

- A neural network is a type of machine learning model that is composed of layers of interconnected "neurons" which process and transmit information.
- Each neuron receives input from other neurons, processes it using a nonlinear activation function, and then transmits the output to other neurons in the next layer. The output of the final layer is the prediction made by the neural network.
- Because all the inputs are densely connected to all outputs, these layers are called Dense layers.
- The output layer can be either single output or multiple output.

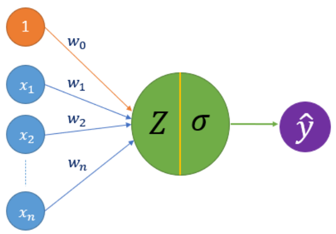

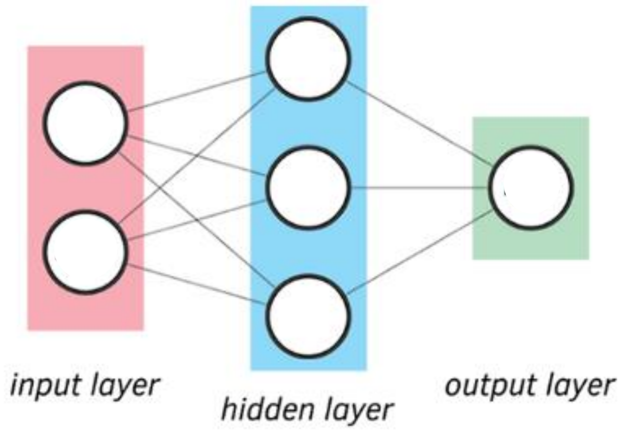

JON M. HUNTSMAN SCHOOL OF BUSINESS
UtahStateUniversity

Pedram Jahangiry

Building a Deep Neural Network

- Deep neural networks are neural networks with a large number of layers, typically consisting of multiple hidden layers.
- Deep neural networks can learn and model very complex patterns in data.
- DNNs have been successful in a wide range of applications, including image and speech recognition, natural language processing, and machine translation.

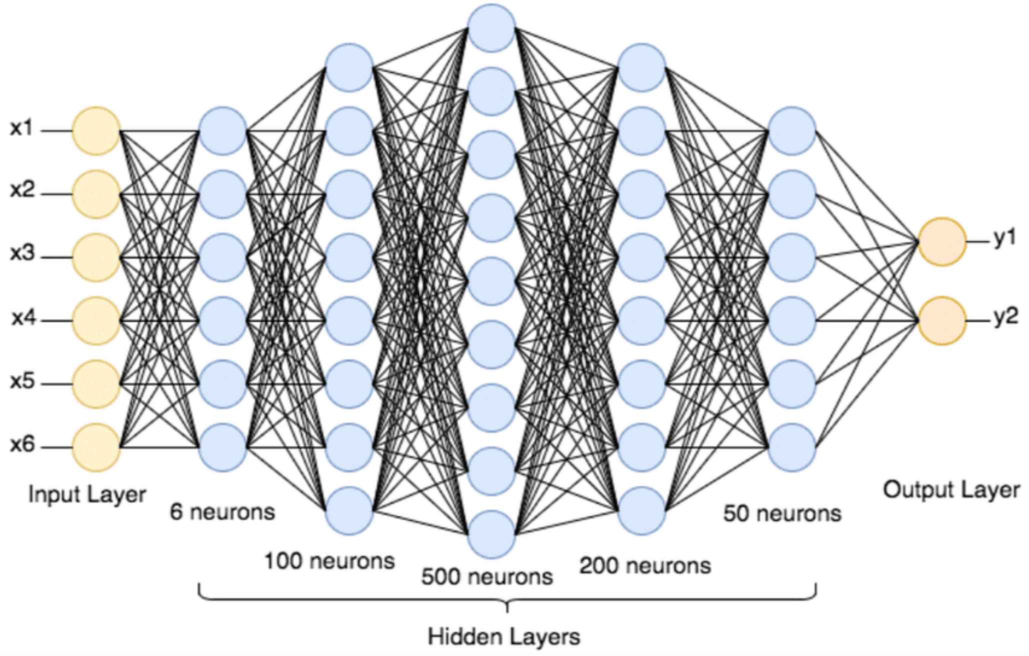


```python
import tensorflow as tf

# Define the model architecture
model=tf.keras.Sequential([
    tf.keras.layers.Dense(6  , activation='relu'),
    tf.keras.layers.Dense(100, activation='relu'),
    tf.keras.layers.Dense(500, activation='relu'),
    tf.keras.layers.Dense(200, activation='relu'),
    tf.keras.layers.Dense(50 , activation='relu'),
    tf.keras.layers.Dense(2  , activation='sigmoid'),
])
```

Jon M. Huntsman School of Business
Utah State University

Pedram Jahangiry

Training a Neural Network, Backpropagation

- To train a neural network, we present it with many examples and adjust the weights and biases of the connections between neurons so that the network can accurately predict the output for each example and minimize the loss function.
- This process is known as backpropagation, and it is done using an optimization algorithm such as stochastic gradient descent.
- The loss function quantifies the distance between actuals and predictions. It provides feedback to the NN.
- Examples: MSE, Binary Cross Entropy, (Sparse) Categorical Cross Entropy, …
- For the full list, visit https://keras.io/api/losses/#available-losses

Jon M. Huntsman School of Business
Utah State University

Pedram Jahangiry

Backpropagation

- Backpropagation is a supervised learning algorithm that adjusts the weights and biases of the connections between neurons in the network to minimize the loss function.

$$\boldsymbol{W}^* = \underset{\boldsymbol{W}}{\operatorname{argmin}} \frac{1}{n}\sum_{i=1}^{n} \mathcal{L}\left(f\left(x^{(i)}; \boldsymbol{W}\right), y^{(i)}\right)$$

- Steps:

1) Feed the input data through the neural network to compute the output.
2) Calculate the loss or error between the predicted output and the true output.
3) Propagate the error back through the network using the chain rule of calculus to calculate the gradient of the loss function with respect to the weights and biases.
4) Update the weights and biases using the gradient descent algorithm.
5) Repeat the process until the loss reaches a satisfactory level or a predetermined iterations.

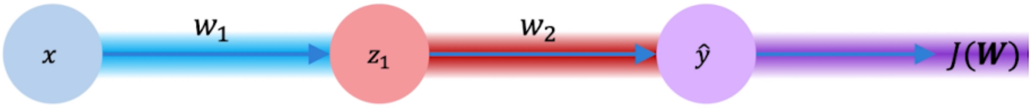

$$\frac{\partial J(\boldsymbol{W})}{\partial w_2} = \frac{\partial J(\boldsymbol{W})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_2}$$

$$\frac{\partial J(\boldsymbol{W})}{\partial w_1} = \frac{\partial J(\boldsymbol{W})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1} * \frac{\partial z_1}{\partial w_1}$$

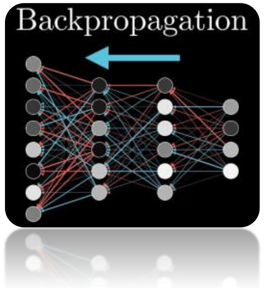

Jon M. Huntsman School of Business
Utah State University

Pedram Jahangiry

Loss optimization

$$\boldsymbol{W}^* = \underset{\boldsymbol{W}}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^{n} \mathcal{L}\left(f\left(x^{(i)}; \boldsymbol{W}\right), y^{(i)}\right)$$

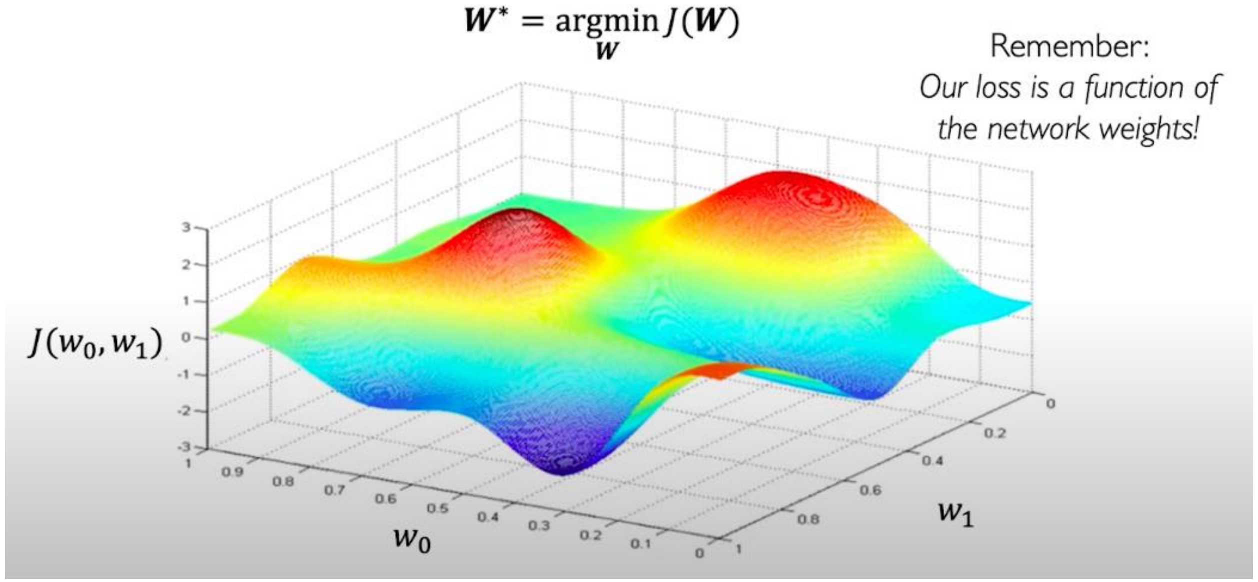

Gradient Descent

1. Initialize weights randomly $\sim\mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\boldsymbol{W})}{\partial \boldsymbol{W}}$
4. Update weights, $\boldsymbol{W} \leftarrow \boldsymbol{W} - \eta \frac{\partial J(\boldsymbol{W})}{\partial \boldsymbol{W}}$
5. Return weights

Loss optimization

$$W^* = \underset{W}{\operatorname{argmin}} \frac{1}{n}\sum_{i=1}^{n} \mathcal{L}\left(f\left(x^{(i)}; W\right), y^{(i)}\right)$$

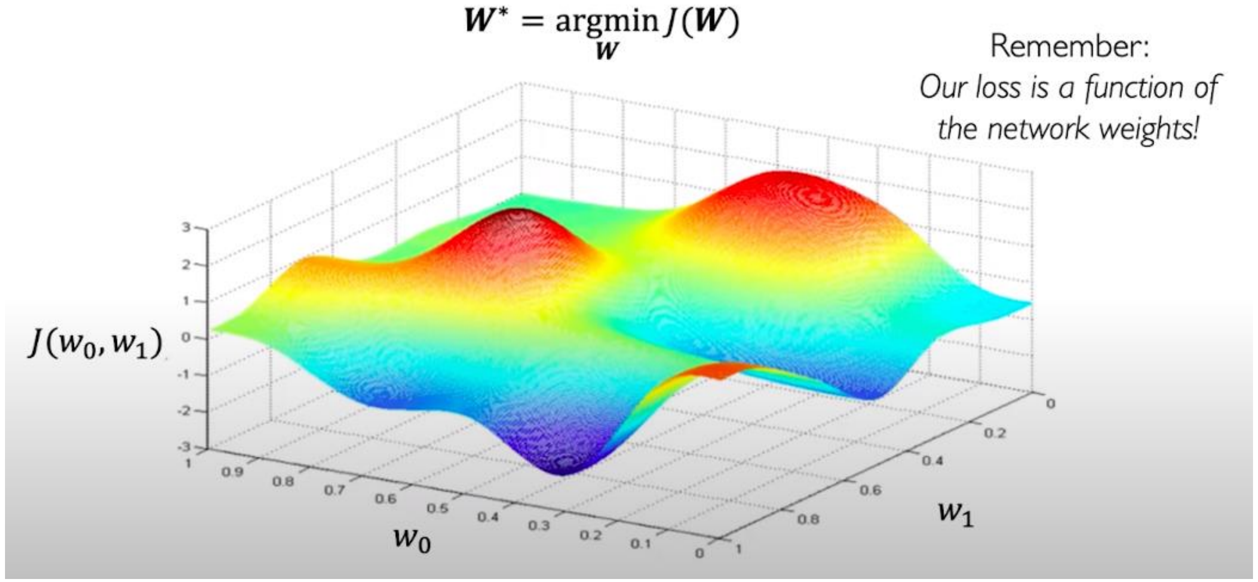

Stochastic GD

1. Initialize weights randomly $\sim\mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J_i(W)}{\partial W}$
4. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

Jon M. Huntsman School of Business
Utah State University

Pedram Jahangiry

Loss optimization

$$W^* = \underset{W}{\operatorname{argmin}} \frac{1}{n}\sum_{i=1}^{n} \mathcal{L}\left(f\left(x^{(i)}; W\right), y^{(i)}\right)$$

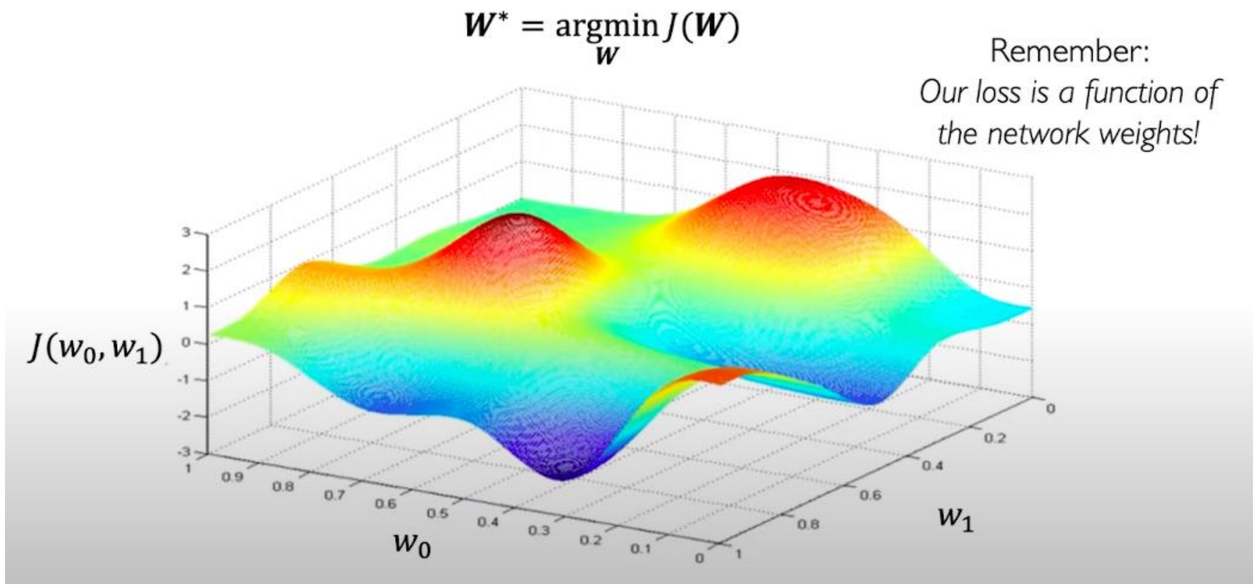

Mini-batch GD

1. Initialize weights randomly $\sim\mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(W)}{\partial W} = \frac{1}{B}\sum_{k=1}^{B} \frac{\partial J_k(W)}{\partial W}$
4. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

Jon M. Huntsman School of Business
UtahStateUniversity

Pedram Jahangiry

DNN loss functions can be difficult to optimize!

- *Visualizing the loss landscape of neural nets*, Li et all, 2018

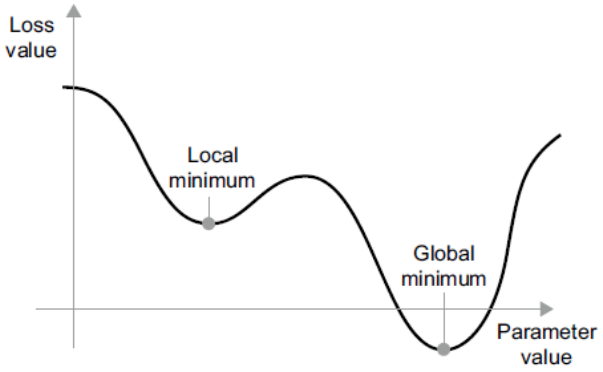

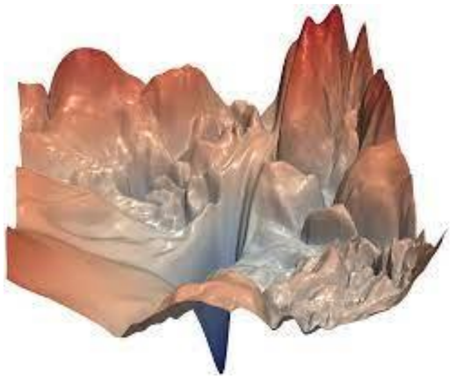

- Solution: Designing an adaptive learning rate that can adapt to the loss landscape. Rather than just looking at the current gradient, take into account the previous weight updates. This is called, momentum!
- Examples: Adam, Adadelta, Adagrad, RMSProp!

Gradient Descent Algorithms

- Available optimizers in Keras:
  - SGD
  - RMSprop
  - Adam
  - AdamW
  - Adadelta
  - Adagrad
  - Adamax
  - Adafactor
  - Nadam
  - Ftrl

Keras

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential()
model.add(layers.Dense(64, kernel_initializer='uniform', input_shape=(10,)))
model.add(layers.Activation('softmax'))

opt = keras.optimizers.Adam(learning_rate=0.01)
model.compile(loss='categorical_crossentropy', optimizer=opt)
```

```
lr_schedule = keras.optimizers.schedules.ExponentialDecay(
    initial_learning_rate=1e-2,
    decay_steps=10000,
    decay_rate=0.9)
optimizer = keras.optimizers.SGD(learning_rate=lr_schedule)
```

- Which optimizer? There is no “right” answer.

Module 4 - Part II
Deep Neural Networks
Regularization

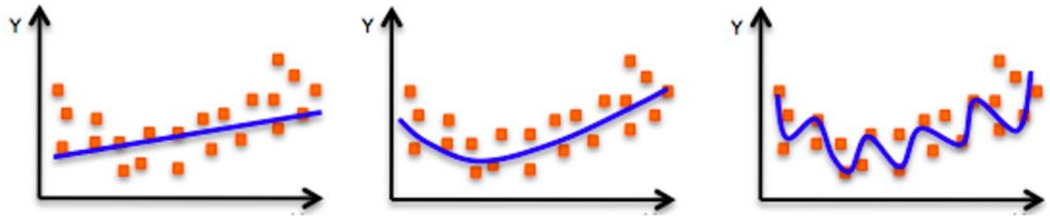

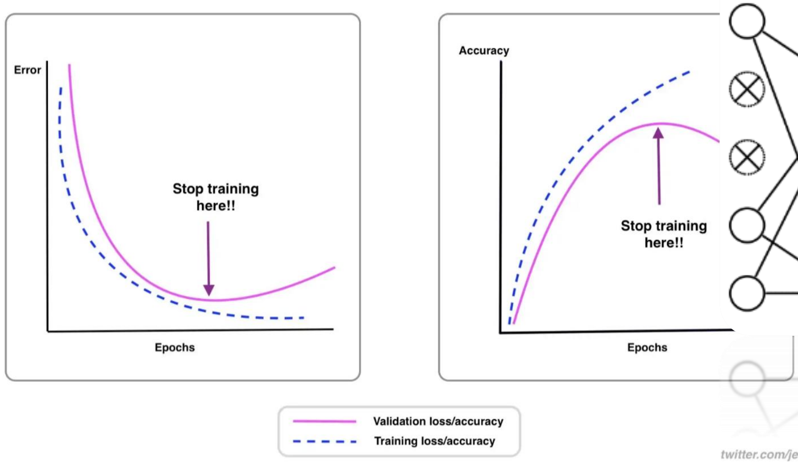

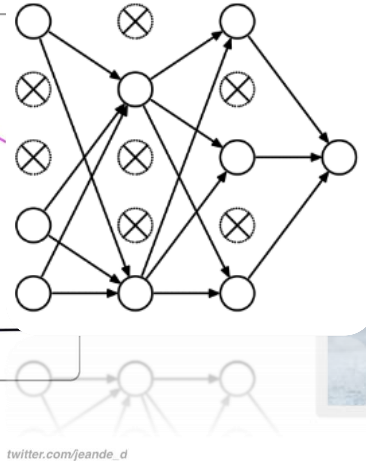

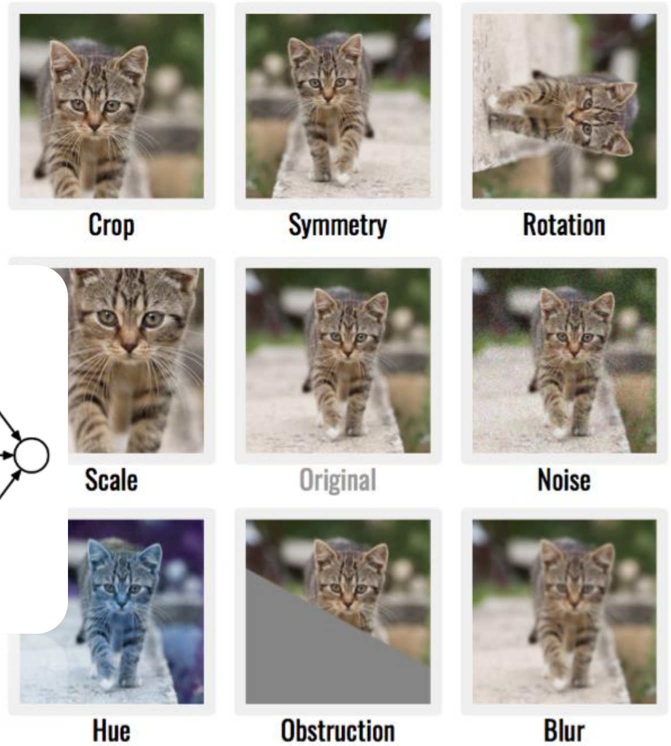

How to handle overfitting in DNN?

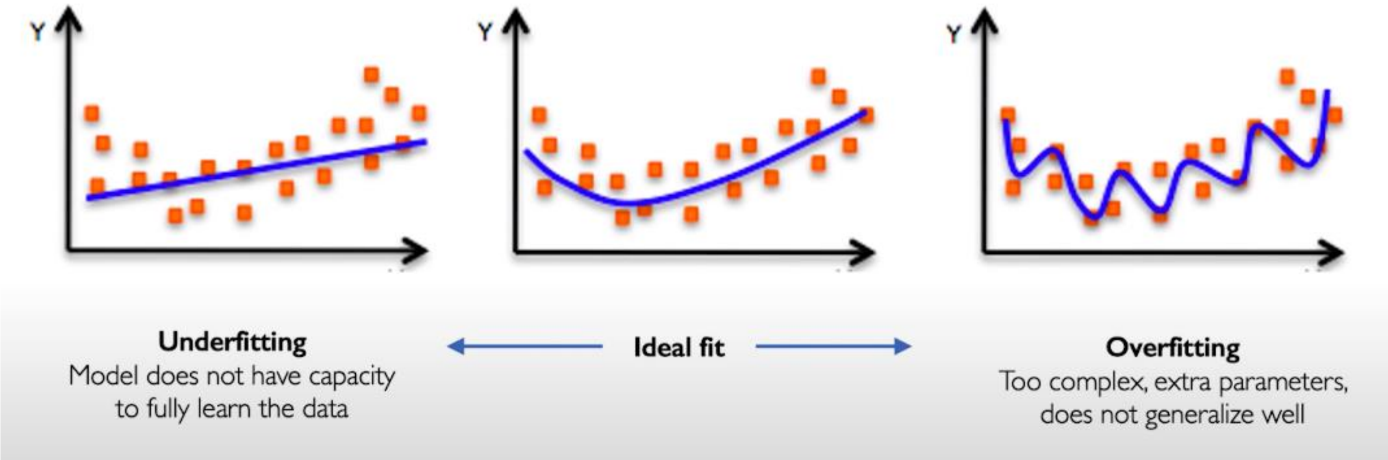

- Regularization refers to a set of techniques that are used to prevent overfitting by discouraging complex models.
- Regularization improve generalization of our model on unseen data.
- Methods include, L1, L2, drop out, Early stopping and Data augmentation.

Deep Learning Regularization techniques

L1 regularization: This technique adds a penalty term to the objective function that is proportional to the absolute value of the model weights. This results in a sparse model, with many weights being set to zero.

L2 regularization: This technique adds a penalty term to the objective function that is proportional to the square of the model weights. This results in a model with small, non-zero weights.

Dropout: This technique randomly sets a fraction of the model weights to zero during training, which helps to prevent overfitting

Early Stopping: This technique involves training the model until the performance on a validation set begins to degrade, and then stopping the training at that point. This helps to prevent the model from continuing to fit the training data too closely.

Data Augmentation: This technique involves generating additional training examples by applying random transformations to the existing training data. This can help to prevent overfitting by providing the model with a more diverse training set.

Dropout Regularization

- Dropout is one of the most popular regularization techniques in deep learning.
- At each training iteration, dropout randomly chooses different nodes to ignore
- This prevents the network from relying on any single neuron.

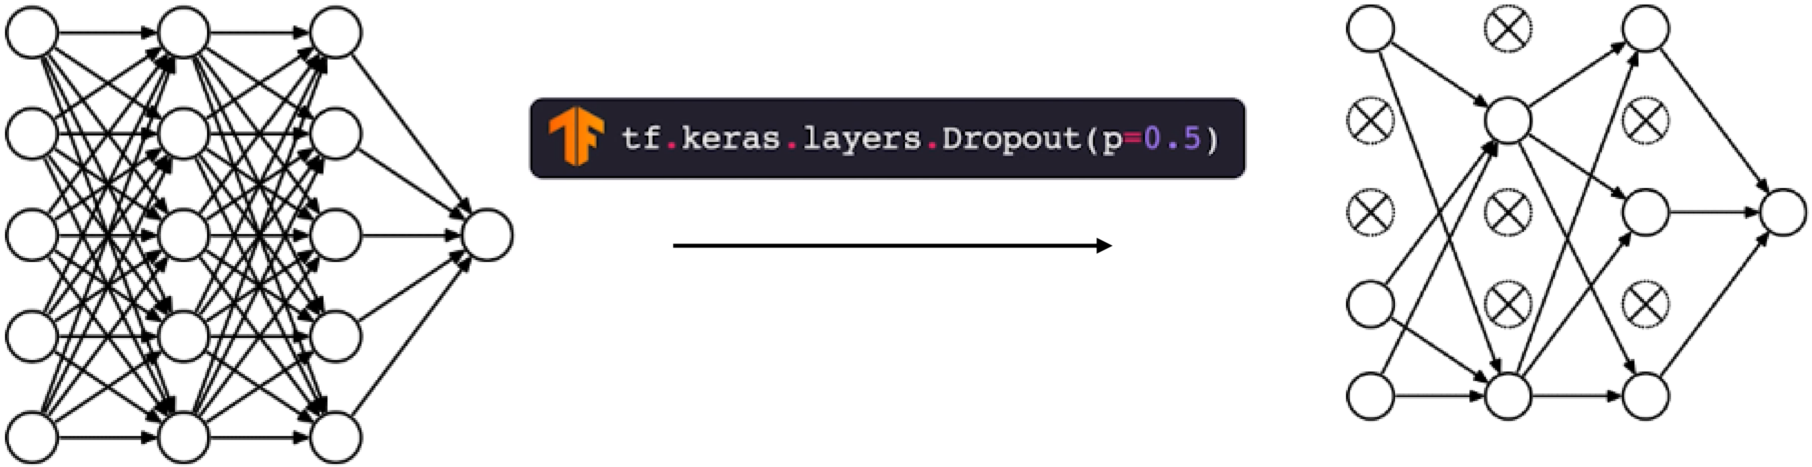

Early Stopping

- Stop training when the performance on a validation set begins to degrade!

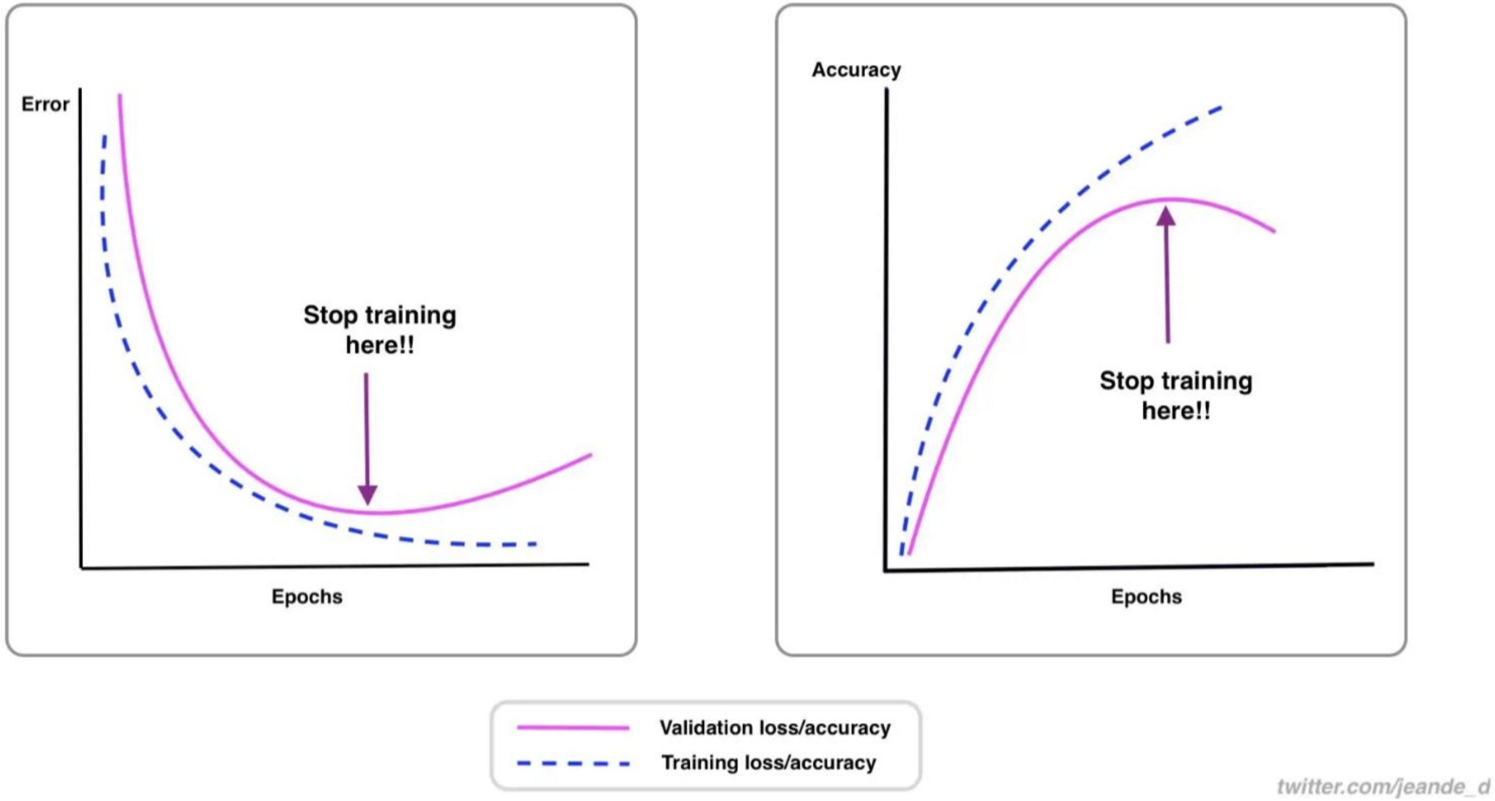

Data Augmentation

- Generating additional training examples by applying random transformations to the existing training data
- The goal of data augmentation is to improve the generalizability and robustness of machine learning models by providing them with more diverse training data.

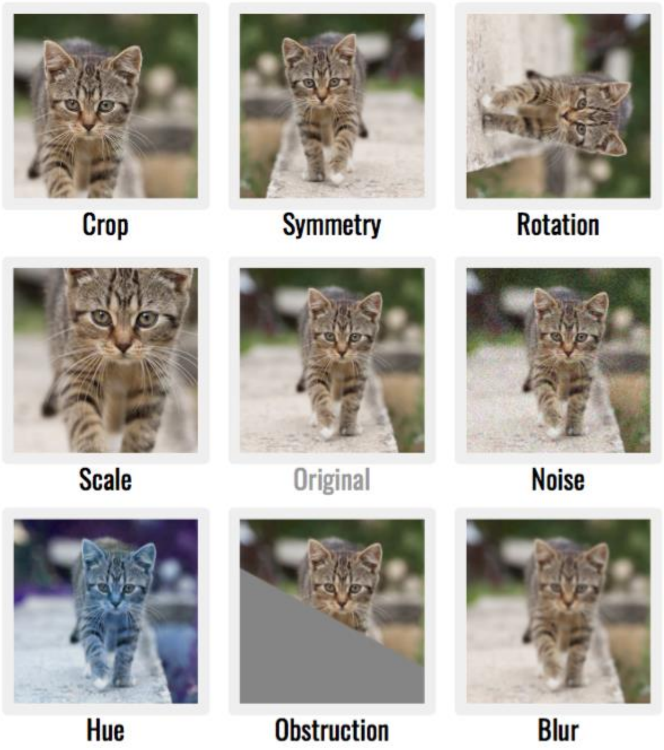

Neural Network Playground (TensorFlow)

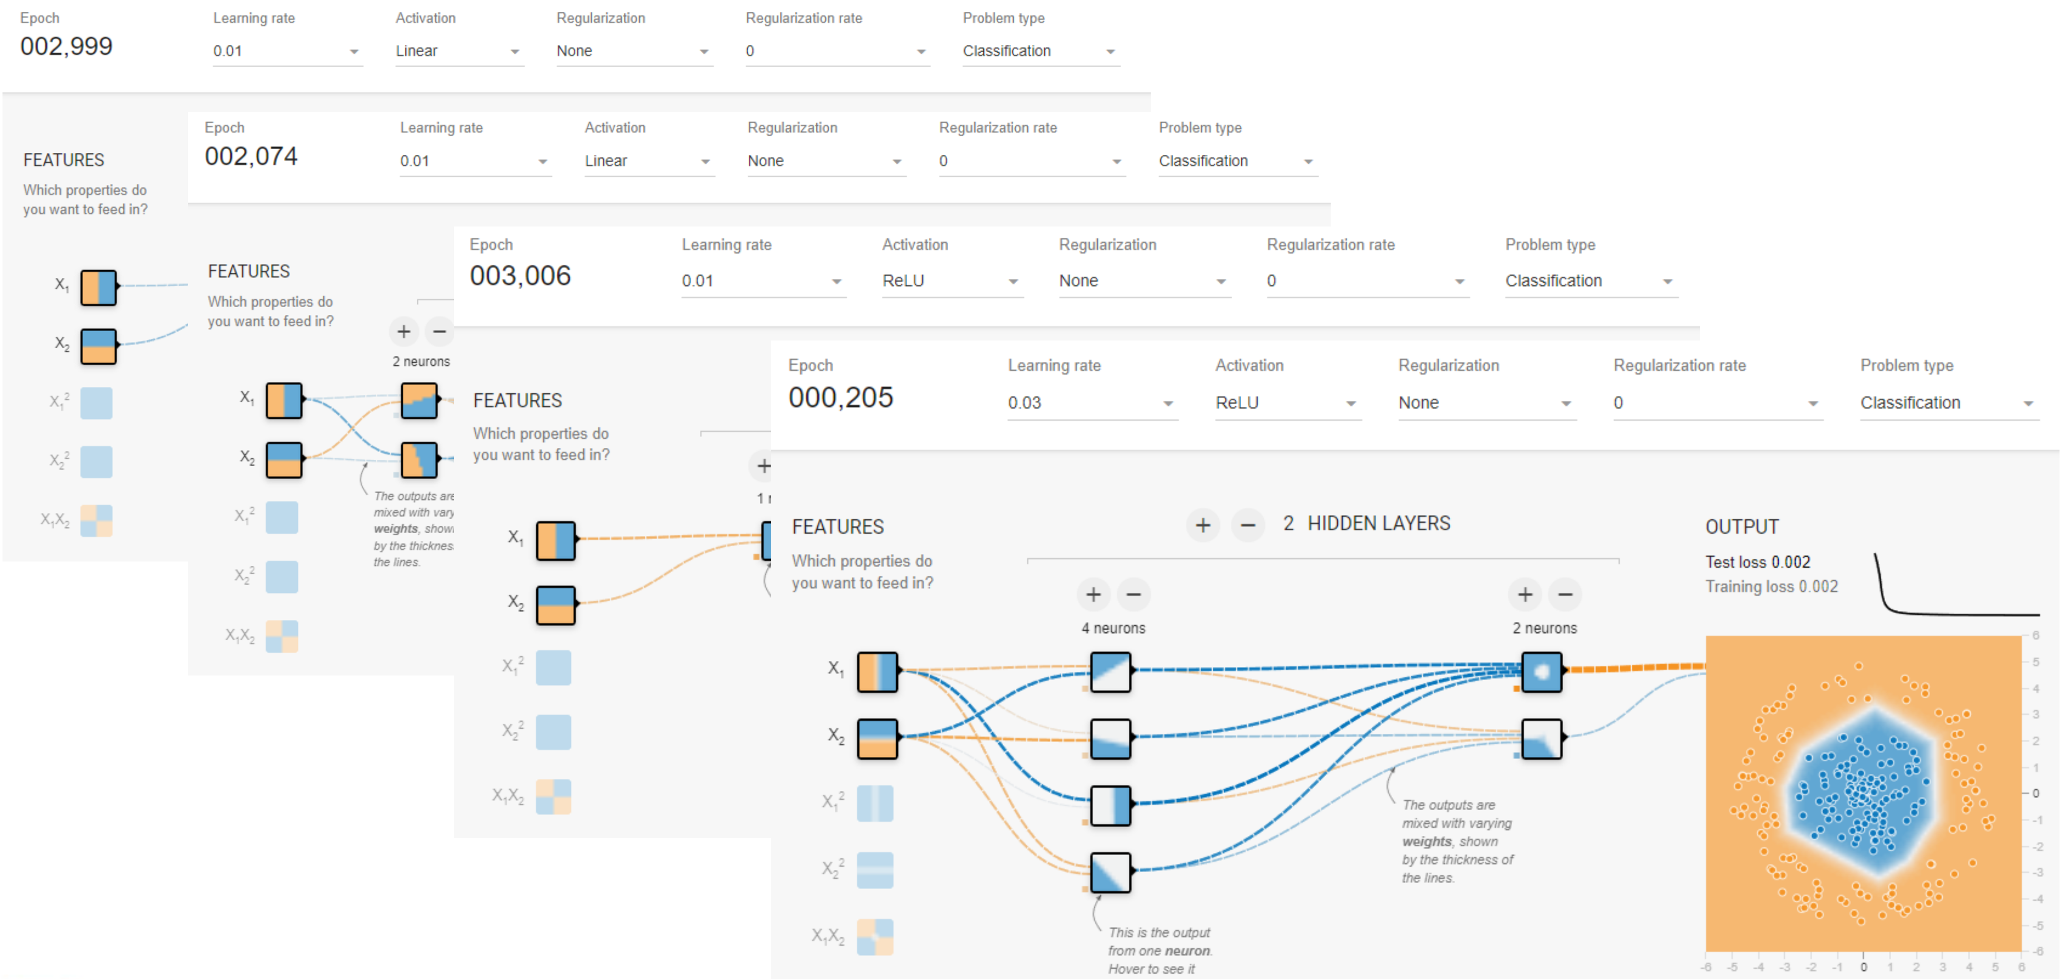

Jon M. Huntsman School of Business
UtahStateUniversity

Pedram Jahangiry

What is a Tensor?

- Tensor: A multi-dimensional array (kind of like np.arrays)
- Rank: Number of tensor axes , Size: The total number of items in the tensor

Tensor	tf code	shape	Example
Scalar (Rank-0 tensor)	tf.constant(4)	Shape = ()	4
Vector (Rank-1 tensor)	tf.constant([2.0, 3.0, 4.0])	Shape = (3 ,)	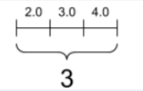
Matrix (Rank-2 tensor)	tf.constant([[1, 2],[3, 4],[5, 6]])	Shape = (3 , 2)	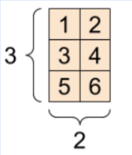 
Rank-3 tensor	tf.constant([[[0, 1, 2, 3, 4], [5, 6, 7, 8, 9]], [[10, 11, 12, 13, 14], [15, 16, 17, 18, 19]], [[20, 21, 22, 23, 24], [25, 26, 27, 28, 29]],])	Shape = (3, 2, 5)	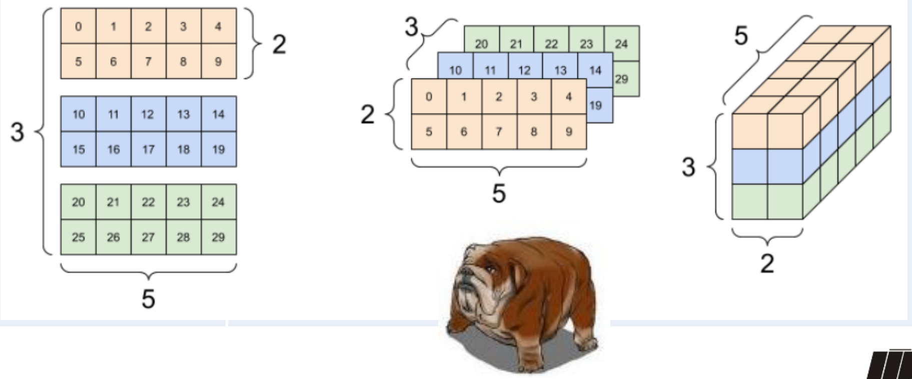

Pedram Jahangiry

https://www.tensorflow.org/guide/tensor

Higher dimension Tensors

```
rank_4_tensor = tf.zeros([3, 2, 4, 5])
```

A rank-4 tensor, shape: [3, 2, 4, 5]

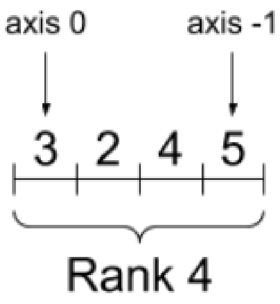

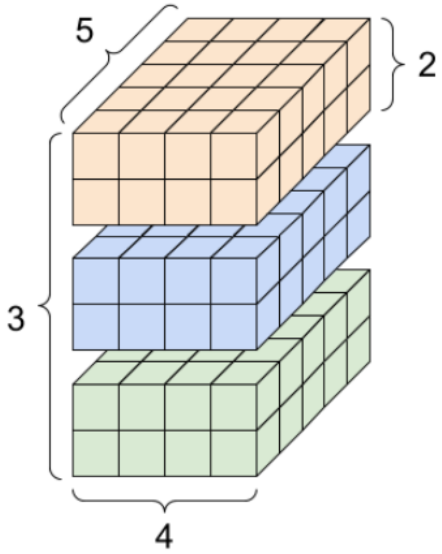


```
print("Type of every element:", rank_4_tensor.dtype)
print("Number of axes:", rank_4_tensor.ndim)
print("Shape of tensor:", rank_4_tensor.shape)
print("Elements along axis 0 of tensor:", rank_4_tensor.shape[0])
print("Elements along the last axis of tensor:", rank_4_tensor.shape[-1])
print("Total number of elements (3*2*4*5): ", tf.size(rank_4_tensor).numpy())
```

```
Type of every element: <dtype: 'float32'>
Number of axes: 4
Shape of tensor: (3, 2, 4, 5)
Elements along axis 0 of tensor: 3
Elements along the last axis of tensor: 5
Total number of elements (3*2*4*5):  120
```

- Often axes are ordered from global to local.
- The first axes is called the batch axis or batch dimension.

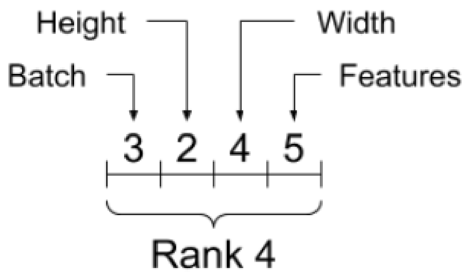

Real-world examples of data tensors

- The data we'll manipulate almost always fall into one of the following categories:

✓ Tabular data: Rank-2 tensors of shape (samples, features), where each sample is a vector of numerical attributes ("features")

✓ Timeseries data or sequence data: Rank-3 tensors of shape (samples, timesteps, features), where each sample is a sequence (of length timesteps) of feature vectors

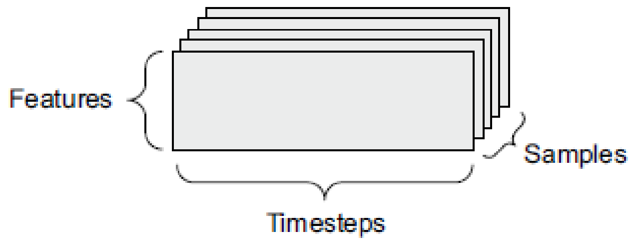

Real-world examples of data tensors

- ✓ Images: Rank-4 tensors of shape (samples, height, width, channels), where each sample is a 2D grid of pixels, and each pixel is represented by a vector of values ("channels")

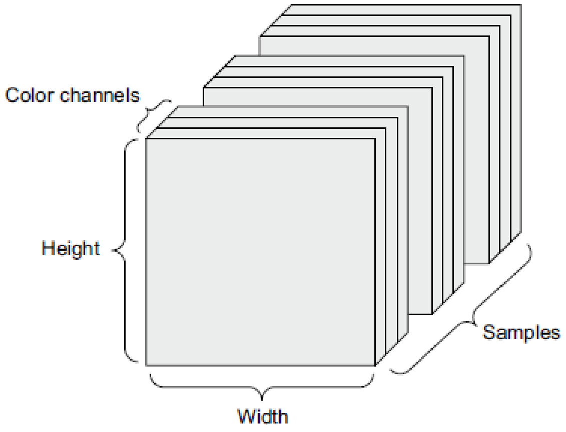

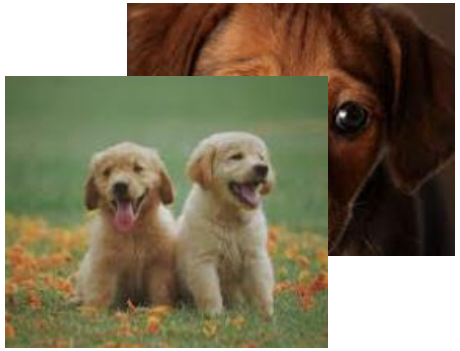

- ✓ Video: Rank-5 tensors of shape (samples, frames, height, width, channels), where each sample is a sequence (of length frames) of images.

Example: What is the tensor shape for a 60-second, 144*256 video clip sampled at 4 frames per second?

Types of Neural Networks

Standard NN

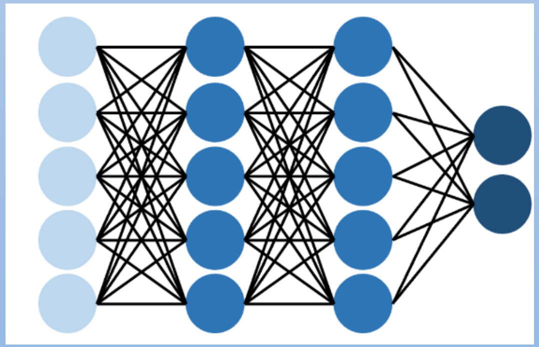

Convolutional NN

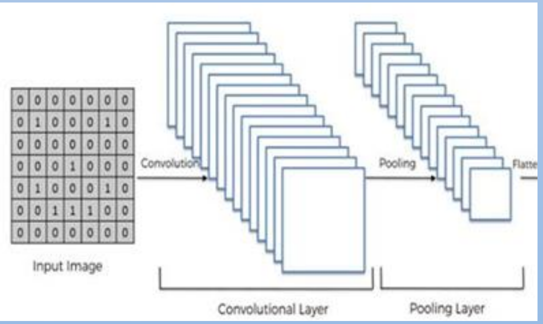

Transformers

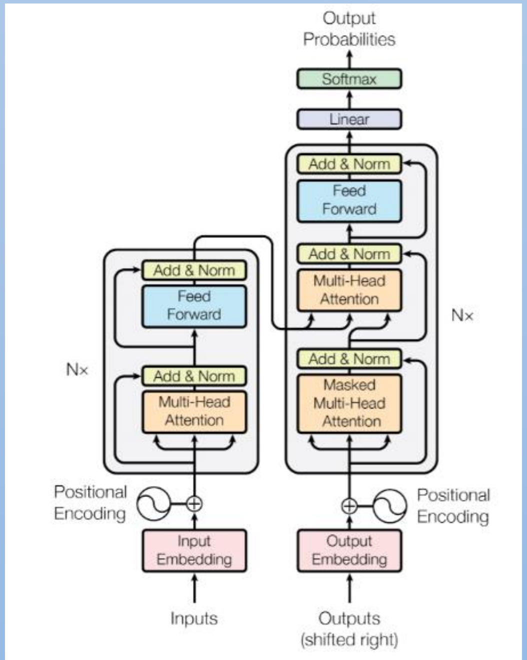

Recurrent NN

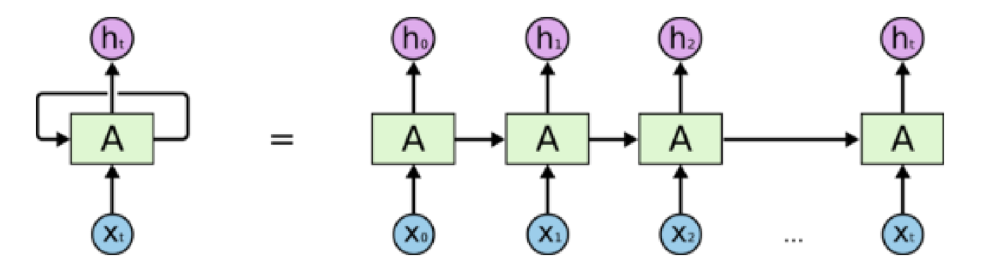

JON M. HUNTSMAN SCHOOL OF BUSINESS
UtahStateUniversity

DNN using 2 lags

```
i = Input(shape=(Tx,))
x = Dense(32, activation='relu')(i)
x = Dense(16, activation='relu')(x)
output = Dense(Ty , activation = 'linear')(x)
model = Model(i, output)
model.compile(loss='mse', optimizer='adam')
```

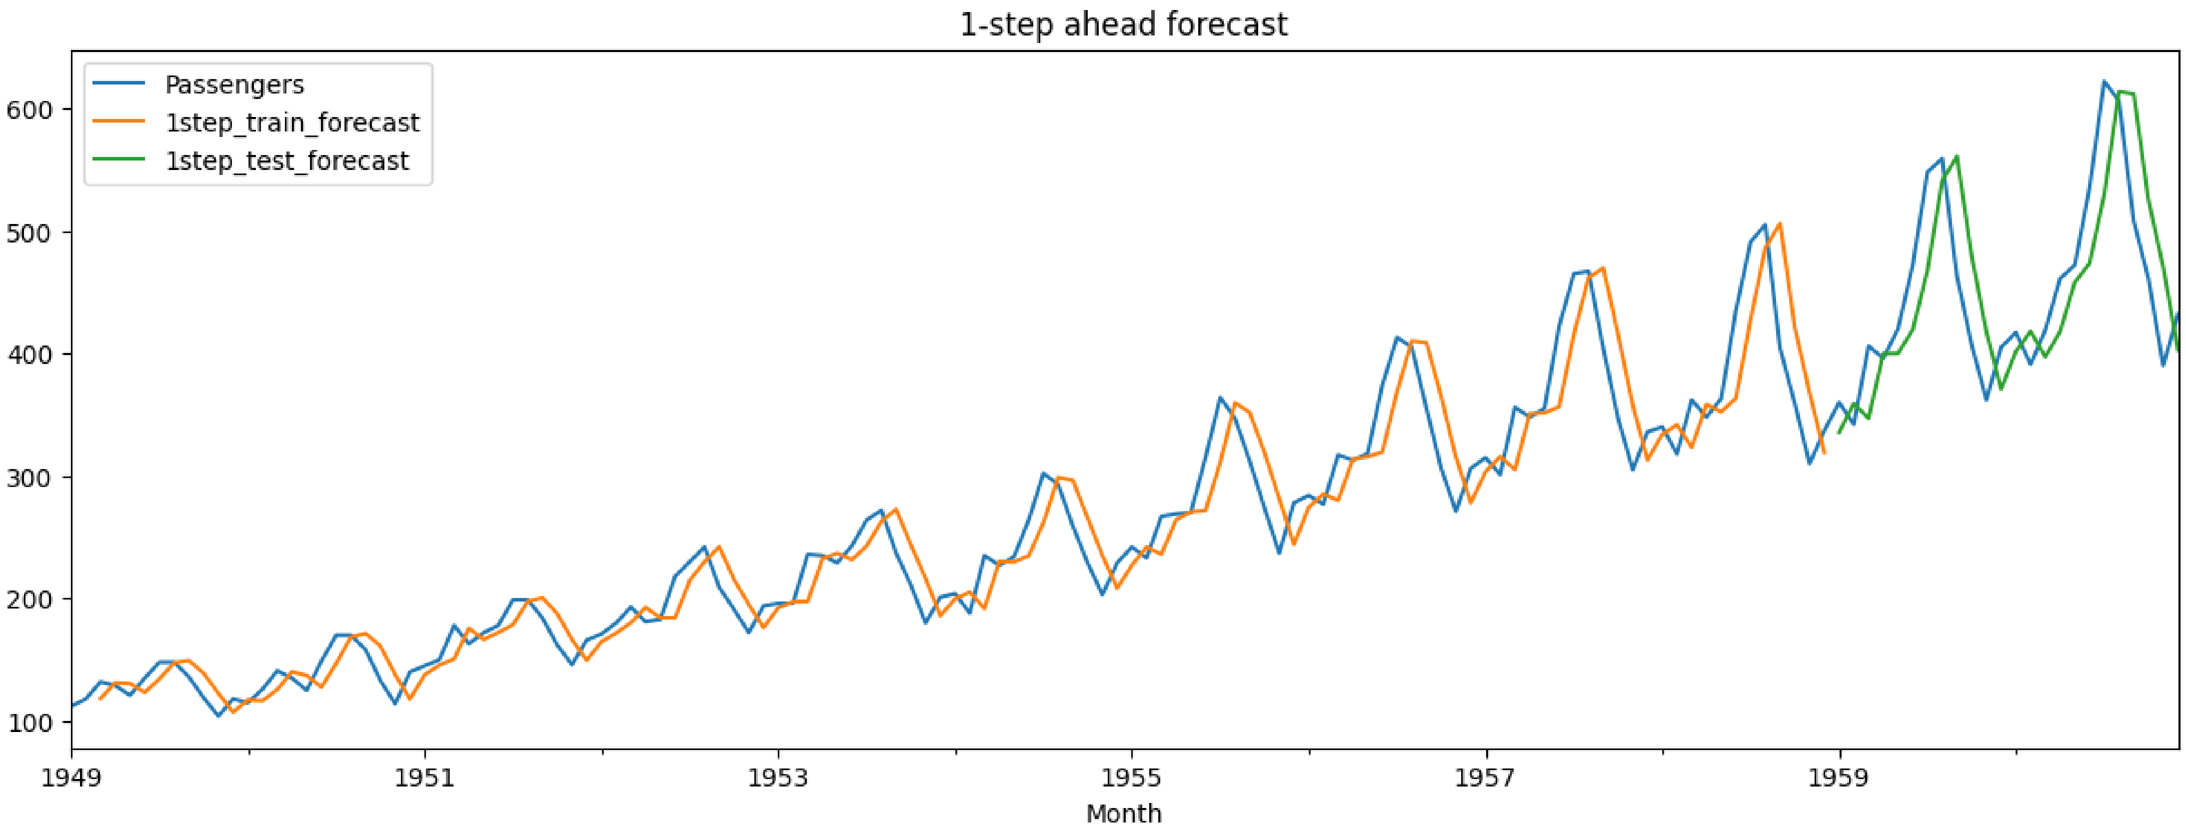

DNN using 2 lags

```
i = Input(shape=(Tx,))
x = Dense(32, activation='relu')(i)
x = Dense(16, activation='relu')(x)
output = Dense(Ty , activation = 'linear')(x)
model = Model(i, output)
model.compile(loss='mse', optimizer='adam')
```

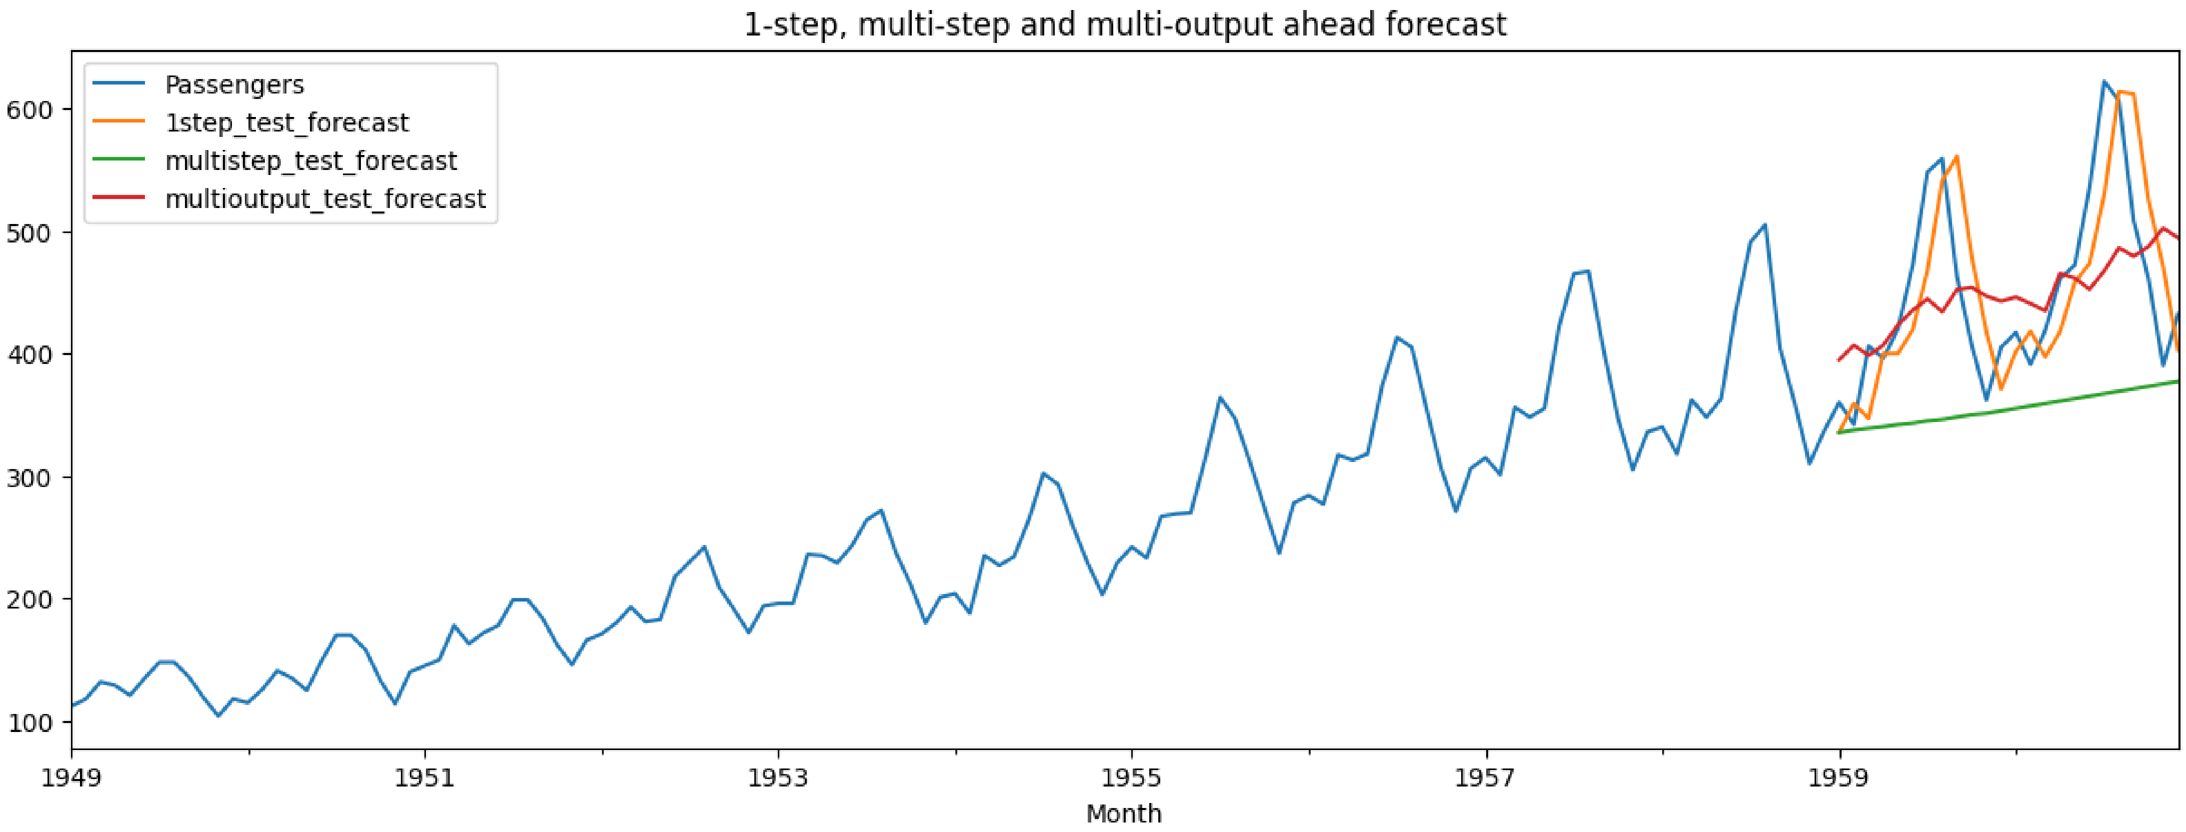

DNN using 12 lags

```
i = Input(shape=(Tx,))
x = Dense(32, activation='relu')(i)
x = Dense(16, activation='relu')(x)
output = Dense(Ty , activation = 'linear')(x)
model = Model(i, output)
model.compile(loss='mse', optimizer='adam')
```

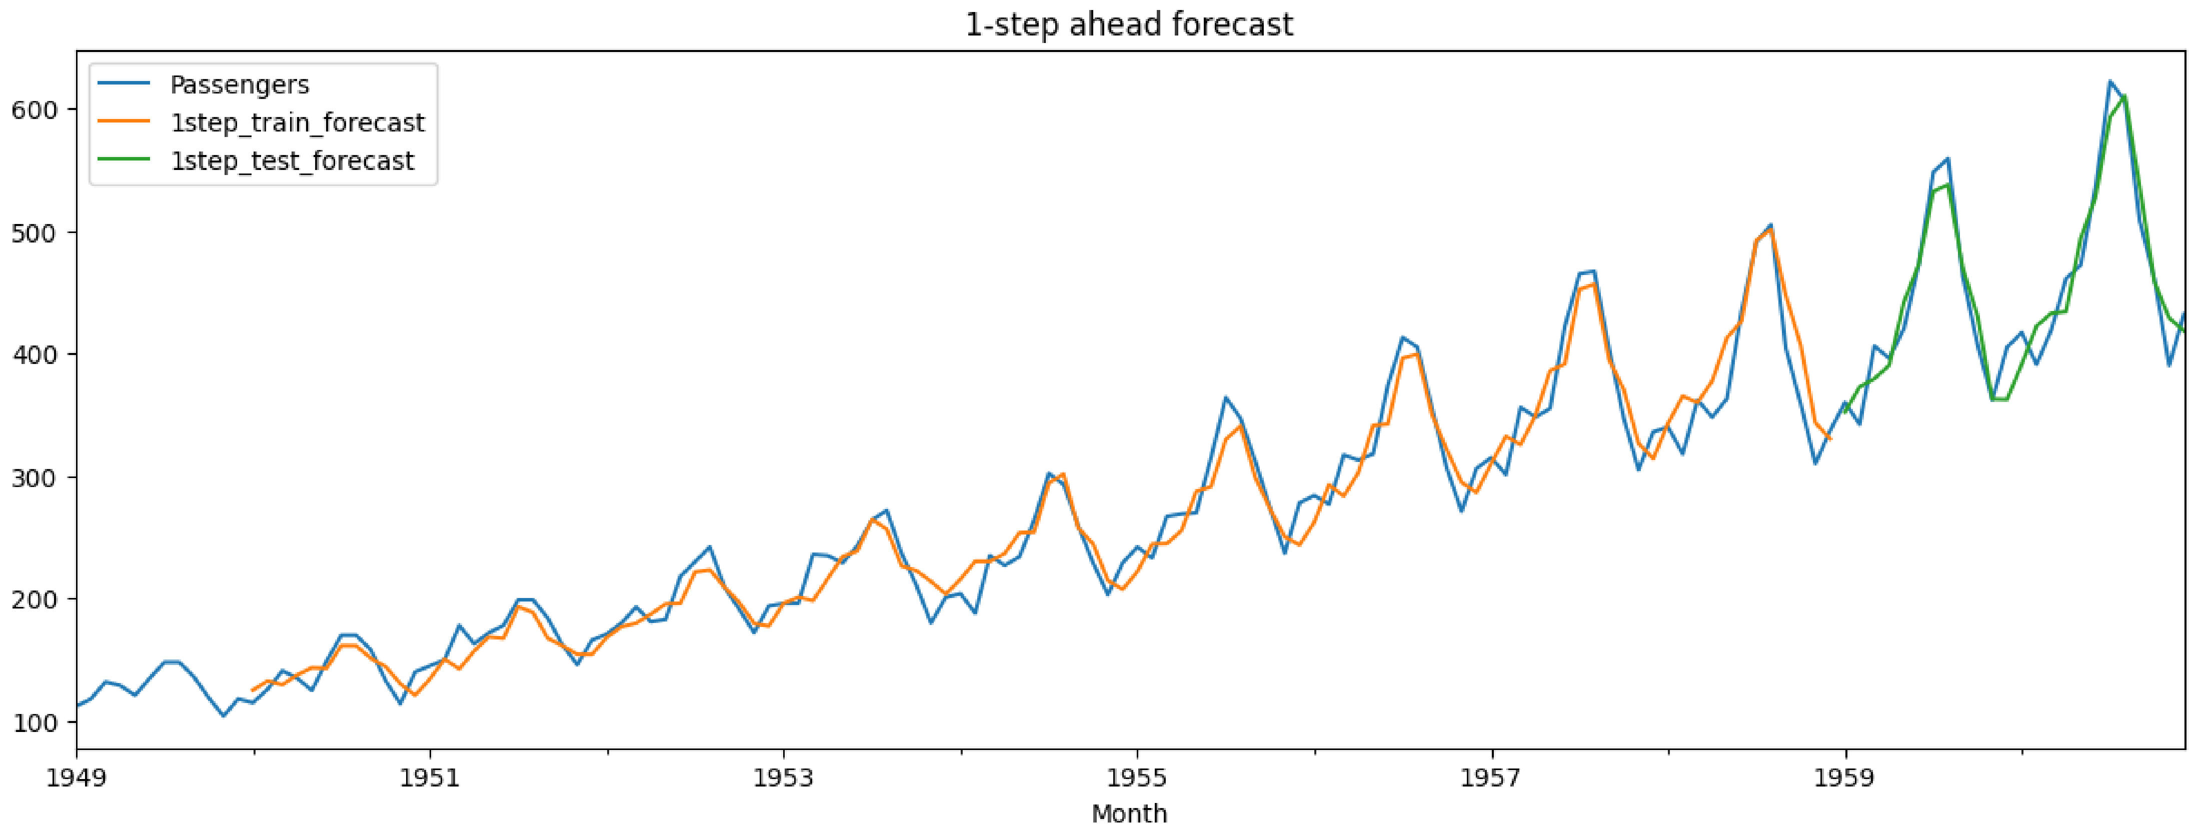

JON M. HUNTSMAN SCHOOL OF BUSINESS
UtahStateUniversity

Pedram Jahangiry

DNN using 12 lags

```python
i = Input(shape=(Tx,))
x = Dense(32, activation='relu')(i)
x = Dense(16, activation='relu')(x)
output = Dense(Ty , activation = 'linear')(x)
model = Model(i, output)
model.compile(loss='mse', optimizer='adam')
```

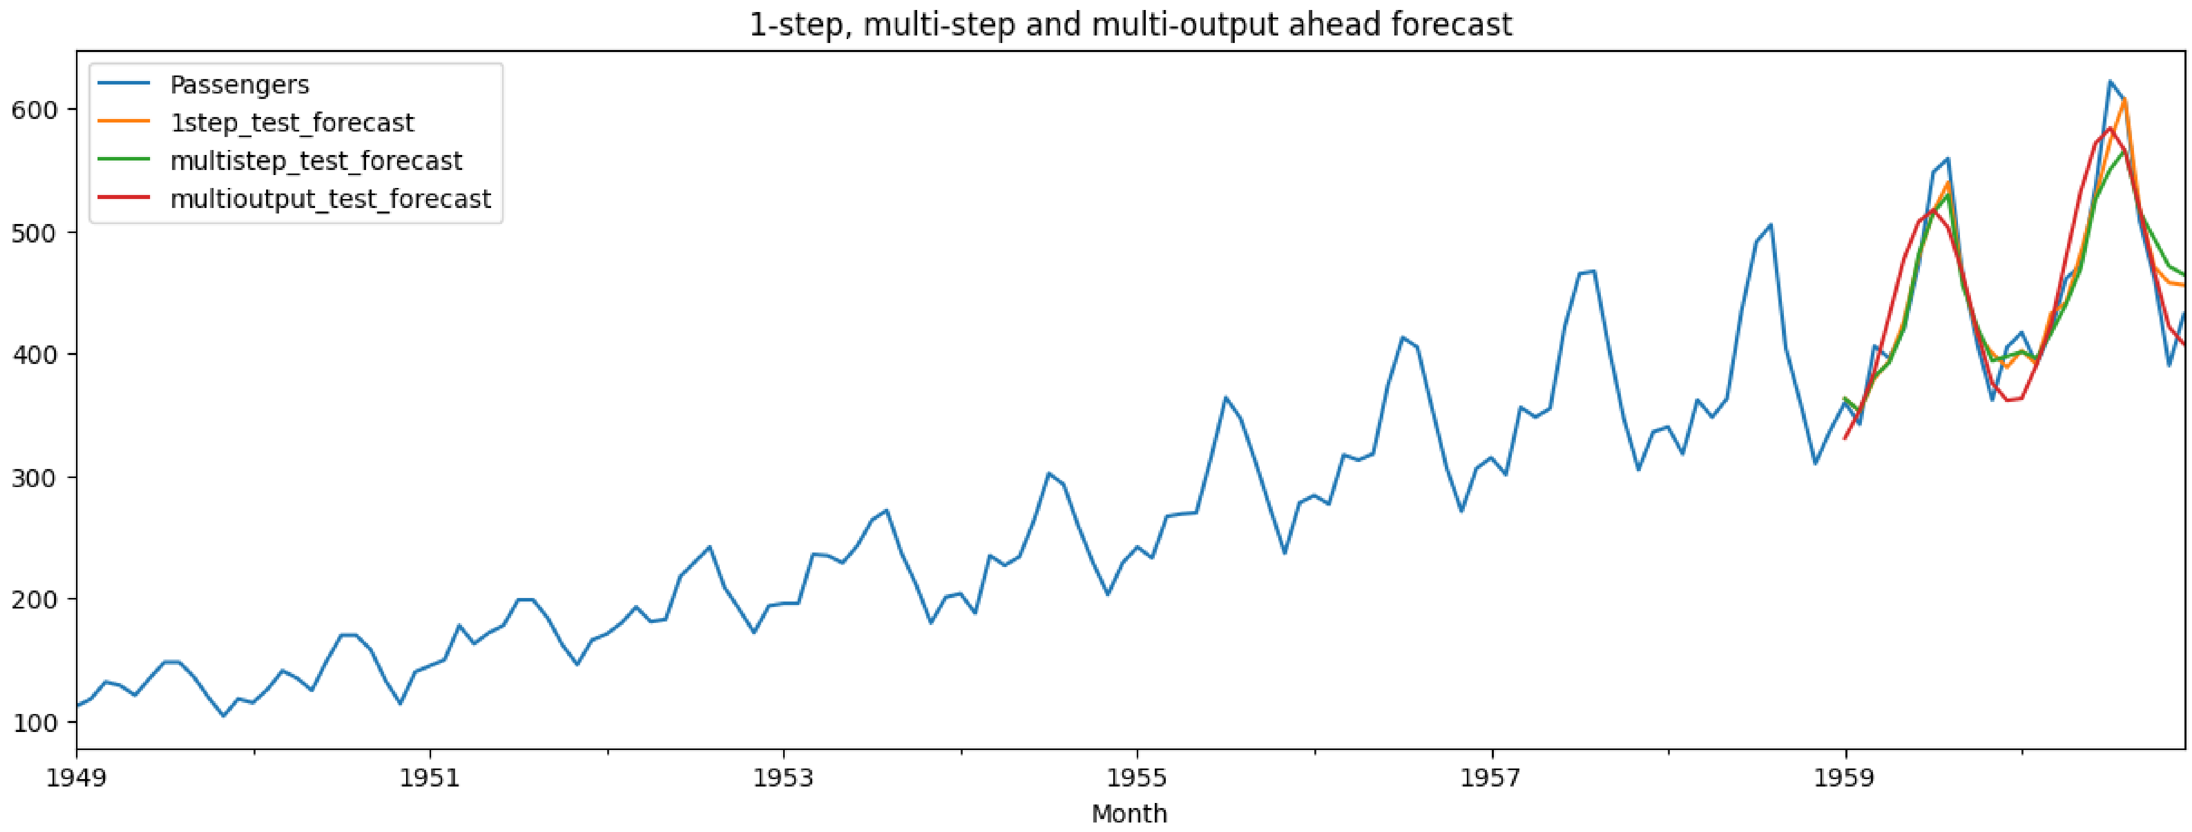

Road map!

- ✓ Module 1- Introduction to Deep Forecasting
- ✓ Module 2- Setting up Deep Forecasting Environment
- ✓ Module 3- Exponential Smoothing
- ✓ Module 4- ARIMA models
- ✓ Module 5- Machine Learning for Time series Forecasting
- ✓ Module 6- Deep Neural Networks
- Module 7- Deep Sequence Modeling (RNN, LSTM)
- Module 8- Transformers (Attention is all you need!)
- Module 9- Prophet and Neural Prophet

Jon M. Huntsman School of Business
Utah State University

Pedram Jahangiry